

Where the Framework Stops

V. The exact form of the obstruction, and the routes that reach it

Abstract

Paper IV collects what the cell coordinates prove about twin pairs. This paper states what stops them, in the form the construction produces on its own rather than in the form sieve theory supplies. **The statement comes first, and the routes that reach it follow.**

Cutting the sieve at depth z with z^3 above the window makes every surviving endpoint either prime or a product of exactly two primes, each with a **unique** responsible line (Theorems 1, 2). Writing C for the surviving cells, R for the composite endpoints inside them and S for the cells with both endpoints composite, the twin count of the window is exactly $T = C - R + S$ (Theorem 4), the deficit $R - C$ can only fall as the cut advances (Theorem 3), and therefore

$$2(R - C) = \sum (-1)^{\Omega(n)}$$

over the endpoints of the surviving cells (Theorem 6). The inequality one wants, $R < C$, is precisely the statement that a **Liouville sum over the sifted set is negative** — equivalently, that more than half of the numbers in (P^2, Q^2) with no small prime factor are prime, for every P and not merely on average. Read as a covariance, the same content puts the missing statement on a correlation between the two rails rather than on any density.

Section 3 shows that no purely sieve-theoretic tool crosses that line. The classical Buchstab upper bound for the composite part equals *exactly twice* the corresponding lower-bound sieve function throughout the relevant range (Theorem 7), so the loss is a clean factor of two uniform in every parameter; the two natural constraints on the cut are incompatible, $R < C$ requiring the sieve variable below $2e^\gamma = 3.5621$ while a positive lower bound in dimension two requires it above $\beta_2 = 4.2664$; switching cannot repair it; and five routes through the line geometry itself are followed in one explicit window and closed by measurement. A closing subsection records the shape of the fallacy the identity is written to avoid, with the counterexample.

Section 4 records, in outline, the routes the framework produced before the identity was written — capacity, resonance and parity separately; the closing budget, which misses by a conditional 0.73% in the wrong direction; a deterministic local constraint that is real but not binding; and the exact inheritance laws of [IV, §5], whose constraint binds one line at a time and dissolves in aggregate. Each arrives at the same place. Two test cases whose answers are known independently — Jacobsthal’s function and almost-primes between squares — then check that the framework reproduces the right shape and supplies no bound (§8).

The parity problem is not an obstacle imported from sieve theory; it is what the construction reduces to on its own. No progress toward the twin-prime conjecture is claimed, and no new bound.

How to read the claims in this paper. Statements set as Theorems, Propositions and Corollaries are proved, and the proofs are given. Everything else falls into two kinds, and we try to keep them apart. A *measurement* is a computation over a stated finite range; it is labelled with that range, and it supports a claim about that range only. A *reading* is our own judgement about what a measurement or a proof appears to mean — where a route seems to lead, why an attempt seems to

fail — and we mark it as ours rather than stating it as established. That distinction matters more here than in Paper IV, because most of this paper is an account of routes that did not work, and the reasons we give for their failing are in several places readings and not proofs.

Keywords: parity obstruction, Buchstab function, Liouville function, Jacobsthal’s function, almost-primes between squares.

MSC 2020: 11N35, 11N05, 11N36.

Summary of the main results

Grouped by section. Each line is informal; the precise statement and its proof are in the section named on the right. Entries in italics are measurements or discussions rather than theorems, and one proposition is labelled measured.

§2 — The obstruction, derived from inside — the core of the paper

Result	What it says	Section
Theorem 1	The depth cut. With z^3 above the window, every surviving endpoint is prime or a product of exactly two primes above z .	§2.1
Theorem 2	Unique owner. A surviving composite lies on the line of its smaller factor only, so it is owed exactly one strike.	§2.1
Theorem 3	Monotone deficit. Advancing the cut changes $R - C$ by $-\sum(k - 1)$, so it never rises and a crossing never reverses. At the final cut $R - C = -T$ — hence the crossing is <i>equivalent</i> to a twin existing, and only the non-reversal is new.	§2.1
Theorem 4	The exact twin count $T = C - R + S$, with no error term.	§2.2
Theorem 5	The pair-overlap bound $T \geq C - R + \frac{2}{3}A + \frac{1}{9}B$, both coefficients extremal. The one place in this work where overlap is a <i>resource</i> and the estimates run the other way.	§2.2
Theorem 6	The parity identity $2(R - C) = \sum(-1)^{\Omega(n)}$ over the surviving endpoints. The inequality $R < C$ is the statement that a Liouville sum over the sifted set is negative.	§2.3
<i>The covariance form</i>	The same content as $T/C = (1 - u)(1 - v) + \kappa$ with κ the covariance of the two composite indicators; twins vanish exactly at maximal anti-correlation, and κ is measured at +0.00005 of the Frechet floor over 793 sectors.	§2.3

§3 — Why an ordinary sieve cannot cross it

Result	What it says	Section
Theorem 7	The classical Buchstab upper bound for the composite part is exactly twice the corresponding lower bound: $I(s) = 2f_1(s)$ on $2 \leq s \leq 4$.	§3.1
Corollary 1	Hence the naive balance is $-f_1(s)$ identically: the loss is a factor of two, not a discrepancy to be tightened away.	§3.1
<i>Closed routes</i>	The two constraints on the cut are incompatible (§3.2); switching the cut cannot repair it (§3.3); five routes through the line geometry, each closed by measurement in one explicit window (§3.4, App. C.1); a control showing the resources are ample, so only the forced residues obstruct (§3.5); and the trap the section exists to avoid, with its counterexample (§3.6).	§3.2–3.6

§4 — The routes that reach the same wall (*in outline; the full accounts are in Appendix C.2*)

Result	What it says	Section
<i>Discussion</i>	Capacity, resonance and parity each fail on their own.	§4, App. C.2
<i>Discussion</i>	The closing budget is decomposed into internal quantities and misses by a conditional 0.73%.	App. C.2.1
Proposition 1	In any contiguous block of length H , each residue of a period- q ruler occurs the stated number of times — a deterministic local constraint, and its weakness.	App. C.2.2
<i>Discussion</i>	Why the three inheritance laws together still do not close the problem.	App. C.2.3

§8 — Scope: two test cases, and the narrowest form of the requirement

Result	What it says	Section
Proposition 2	<i>Measured, not proved.</i> Saturation of extremal runs in the first test case, Jacobsthal’s function — and §8.4 explains why it yields no bound.	§8.3
<i>Second test case</i>	Almost-primes between squares, with the correlation of the survivor count and the order of its sum.	§8.5–8.6
<i>The narrowest form</i>	The requirement reduced as far as the construction can take it — and why it is still the wall.	§8.7

§9 — Conclusion

Result	What it says	Section
<i>What is required</i>	What an external ingredient would have to supply; why this stopping point is the right one; that this is not a new sieve; the limitations; and what the framework does not presently supply.	§9

1. Setting

We use Papers 0–IV as follows and import nothing else.

- [0] supplies the window combinatorics: the increments W_j of $\lfloor 2j^2/p \rfloor$, their uniform bound and their exact histogram.
- [I] supplies the coordinates: the line $L_p(k) = p(p + 2k)$ beginning at p^2 , the grid L_3 , the cells $C_b = (6b - 1, 6b + 1)$.
- [II] supplies the exact cycle laws: the four-state census and its refinement by inheritance depth.
- [III] supplies the passage to a short window, and the measurement $T/M \approx 0.80$ at moving depth.
- [IV] supplies everything this paper tests: the gap alphabet, the six exception positions, the closing budget, the clocks, and the sector inheritance laws.

The three words used constantly below, all as defined in [I]: a **line** L_p is the odd multiples of p from p^2 onward; a **cell** is a pair $C_b = (6b - 1, 6b + 1)$, and it is **open**, or **survives** a set of lines, when neither member lies on any of them; the **window** is the interval between two consecutive odd squares, which in cell coordinates is a block of consecutive indices. A twin pair is an open cell whose two members are both prime, and inside (u^2, v^2) sieved by every line up to u the two notions coincide [IV, §3.1].

The order of this paper is not the order in which it was found. Section 2 states the obstruction and §3 shows that no ordinary sieve crosses it; §§4–7 then give the routes that were tried first, each of which arrives at the same place. Throughout, a statement labelled *measured* is a numerical finding with its controls; it is not a theorem and not an asymptotic claim.

2. The obstruction, derived from inside

This section states the obstruction first, because everything else in this paper is an instance of it. Cutting the sieve at a depth that removes every complication except one leaves a single distinction standing, and §3 shows that no purely sieve-theoretic tool crosses it. Sections 4 to 7 then take the four routes the framework itself suggested, one at a time, and show each of them arriving at exactly this point — they are the anatomy of the wall, not a sequence of separate failures, and a reader who takes the identity of §2.3 on trust may stop at the end of §3.

2.1 The depth cut

Fix a window $W \subset (1, U)$ and a cut z with

$$z^3 > U. \quad (2.1)$$

Theorem 1 (depth cut). After every line up to z has been switched on, each surviving endpoint $n \in W$ is either prime or a product of exactly two primes, both exceeding z .

Proof. Every prime factor of a survivor exceeds z , or a line below z would have closed it. Three such factors would give $n > z^3 > U$, which is impossible. ■

Theorem 2 (unique owner). If $n = qr$ with $q < r$ both prime, then n lies on L_q and not on L_r . Hence every surviving composite endpoint has exactly one line responsible for it: its smaller factor.

Proof. $qr \geq q^2$, so $n \in L_q$. And $qr < r^2$, so L_r has not yet begun at n . ■

Together these remove, at one stroke, the two things a sieve at ordinary depth has to track: the multiplicative depth beyond two, and the overlap of several lines on one endpoint. **After the cut there is exactly one kind of composite left, and exactly one line responsible for each.**

Moving the cut, and why the crossing never reverses. Write $D(z) = R(z) - C(z)$ for the deficit at cut z , where $C(z)$ counts surviving cells and $R(z)$ the composite endpoints inside them, equivalently the strikes still owed by the lines above z . Advance the cut to the next prime r , and let A_r be the surviving cells that r closes; for $X \in A_r$ let $k(X) \geq 1$ be the number of lines in $(z, P]$ that would have struck it, r included.

Theorem 3 (monotone deficit). $D(r) = D(z) - \sum_{X \in A_r} (k(X) - 1) \leq D(z)$.

Proof. A surviving cell outside A_r is by definition not struck by r , so it loses none of its owed strikes; hence $R(z) - R(r) = \sum_{X \in A_r} k(X)$, while $C(z) - C(r) = |A_r|$. Subtracting gives the identity, and $k \geq 1$ gives the inequality. ■

Verified. Non-increasing at every one of the 169 successive cuts for the sector of $P = 1009$ and every one of the 304 for $P = 2003$, with no exception.

What drives it down. A cell whose only future striker is r contributes $k - 1 = 0$ and moves D not at all; cells carrying two or more owed strikes are the whole of the decrease. So here — unlike everywhere else in this work — the piling of several lines on one cell pushes the inequality in the direction one wants, because closing such a cell discards more owed strikes than it discards cells.

And the reading that must go with it. Take the cut to its end, $z = P$. No lines remain above it, so $R(P) = 0$, and every surviving cell is a twin, so $C(P) = T$. Hence

$$D(P) = -T,$$

verified exactly: at $P = 1009$ the last cut gives $C = 54$, $R = 0$, $D = -54$ against 54 twins; at $P = 2003$ it gives $D = -205$ against 205. Since D is non-increasing and ends at $-T$, the existence of a cut with $R < C$ is **equivalent** to $T > 0$. The crossing is therefore not evidence for a twin: it is the same statement in another coordinate. What Theorem 3 does buy is that the crossing cannot reverse, so it suffices to establish $R < C$ at a single convenient cut rather than at all of them.

2.2 The exact twin count

Write, for the window after the cut,

C	surviving cells (both members surviving)
R	composite endpoints inside those cells — a cell with both members composite is counted twice
S	cells with both endpoints composite
T	cells with both endpoints prime, i.e. twin pairs

Theorem 4. $T = C - R + S$.

Proof. Classify the surviving cells by state: $C = T + (\text{one member composite}) + S$, while $R = (\text{one member composite}) + 2S$. Subtracting eliminates the middle class. ■

Verification. Exact on every window tested. At $M = 999$ with the core taken to $P = (M+6)^{2/3} \approx 100$: $C = 234$, $R = 175$, $S = 34$, and $C - R + S = 93$, which is the true number of twin pairs in $(999^2, 1005^2)$. Over a sweep of 56 windows with M prime between 101 and 2500 and $h \in \{M/2, M\}$, the identity held without exception, with $R/C \in [0.711, 0.785]$ and $T/C \in [0.370, 0.410]$.

A quantitative model for the three counts, and the step it does not license. Take the cut at the boundary of (2.1), $z = Q^{2/3}$ in a sector (P^2, Q^2) , and put

$$h = \sum_{z < q \leq P} \frac{1}{q} \rightarrow \log \frac{3}{2} = 0.405465\dots,$$

the limit because $\log P / \log z \rightarrow 3/2$. If the owed strikes fell on the surviving cells independently, one would get

$$\frac{R}{C} \rightarrow 2h = 0.810930, \quad \frac{S}{C} \rightarrow h^2 = 0.164402, \quad \frac{T}{C} \rightarrow (1 - h)^2 = 0.353472.$$

Measured over the 816 sectors with $5 \leq P < 6300$: mean $R/C = 0.797$, rising toward $2h$ with P (0.748, 0.786, 0.798, 0.805 by range), and $T/C \approx 0.35$ throughout. Sample sectors: $1009 \rightarrow 1013$ gives $C = 146$, $R = 110$, $S = 18$, $T = 54$; $2003 \rightarrow 2011$ gives 512, 396, 89, 205; $6229 \rightarrow 6247$ gives 2729, 2229, 454, 954.

The model is good and the inference from it is invalid, which is the point of recording it. The temptation is to argue that $2h < 1$ forces $R < C$. It does not: $2h$ is an average over a long range, while a sector is a short interval, and the two are not interchangeable. The counterexample is small and explicit. In the sector $29^2 \rightarrow 31^2$ the cut is $z = 31^{2/3} = 9.87$ and $2h = 0.7145 < 1$, yet

$$C = 8, \quad R = 8, \quad S = 2, \quad T = 2,$$

so $R = C$ exactly, with Theorem 4 still holding as $8 - 8 + 2 = 2$. Over $5 \leq P < 6300$ it is the only sector with $R \geq C$.

And the margin is not generous. Writing the requirement as $R - 2hC < (1 - 2h)C \approx 0.189C$, the fraction of that margin actually consumed has mean 0.064 and median 0.069, but reaches 1.000 at $P = 29$, 0.873 at $P = 809$, 0.782 at $P = 599$, 0.723 at $P = 1487$ and 0.564 at $P = 3539$. The worst case falls with P , slowly. Equivalently, since $R = 2C(1 - \rho)$ with ρ the proportion of surviving endpoints that are prime, the inequality $R < C$ is $\rho > 1/2$: what has to be proved is that more than half of the z -rough numbers in (P^2, Q^2) are prime, for every P and not on average. Any argument that assumes $\rho > 1/2$ to derive $R < C$ has assumed its conclusion.

The overlap as a resource, and a lower bound that does not need the model at all. Everything above tries to force $R < C$. There is a better use of the same data. For a surviving

cell write ℓ and r for the number of lines above the cut that strike its left and its right member. Under (2.1) — indeed under the weaker $z \geq Q^{1/2}$ — no member carries four such factors, since their product would exceed $z^4 \geq Q^2$, so $\ell, r \leq 3$. Put

$$A = \sum \left[\binom{\ell}{2} + \binom{r}{2} \right], \quad B = \sum \ell r,$$

the pairs of lines piling on one member and the pairs splitting across the two.

Theorem 5 (pair-overlap bound). For all integers $0 \leq \ell, r \leq 3$,

$$\mathbf{1}_{\{\ell=r=0\}} \geq 1 - \ell - r + \frac{2}{3} \left[\binom{\ell}{2} + \binom{r}{2} \right] + \frac{1}{9} \ell r,$$

and therefore, summing over the surviving cells,

$$T \geq C - R + \frac{2}{3}A + \frac{1}{9}B.$$

Proof. Sixteen cases, checked directly; equality holds at $(0,0)$, $(0,1)$, $(1,0)$, $(0,3)$, $(3,0)$ and $(3,3)$. ■

The coefficients cannot be improved in this shape. Seeking $T \geq C - R + aA + bB$, the case $(3,0)$ gives $1 - 3 + 3a \leq 0$, so $a \leq 2/3$; taking $a = 2/3$, the case $(3,3)$ gives $1 - 6 + 4 + 9b \leq 0$, so $b \leq 1/9$.

What changes, and it is the only place in this work where it changes. The bound uses the overlap as a **resource**: the more the lines pile up, the larger A and B , and the better the bound. Consequently the estimates one needs run in the opposite direction from everywhere else — an upper bound for R and lower bounds for A and B . In the free model $R/C \sim 2h$ and $A/C \sim B/C \sim h^2$ give $T/C \gtrsim 1 - 2h + \frac{7}{9}h^2$, whose root is $h_c = (9 - 3\sqrt{2})/7 = 0.679623$, i.e. a cut exponent $\alpha_c = e^{-h_c} = 0.50681$ against 0.60653 for $R < C$ alone.

And the model overstates A , which is the honest limit of this. Measured against the correct truncated model $2e_2$ with $e_2 = \frac{1}{2}(h^2 - \sum 1/q^2)$, the ratio $A_{\text{measured}}/A_{\text{model}}$ is 0.18–0.37 at $\alpha = 0.60$, 0.50–0.67 at 0.55, 0.62–0.79 at 0.53 and 0.76–0.90 at 0.50 — never reaching one, and rising only as α falls, because two lines on the *same* member is a tail event that the cap $\ell \leq 3$ truncates. B by contrast matches, at 0.90–1.15 throughout. Using the measured ratios the effective threshold is near $\alpha \approx 0.54$, not 0.507, and direct evaluation agrees: the bound reads +8.6, +9.6, +1.9 at $\alpha = 0.53$ for $P = 1009, 2003, 6229$, and turns negative at 0.52.

The quantity to press on is therefore A , and it is the first in this work with a classical shape of its own: it counts semiprimes q_1q_2 with both factors above z inside a short interval. A lower bound for A , together with an upper bound for R , would close the inequality at a fixed exponent — and neither is supplied here.

2.3 The parity identity

By Theorem 1 every surviving endpoint has $\Omega(n) \in \{1, 2\}$, so $(-1)^{\Omega(n)}$ is -1 on the primes and $+1$ on the semiprimes. Writing P for the number of prime endpoints, one has $P + R = 2C$ and hence $P = 2C - R$, so

$$\sum_{\text{endpoints of surviving cells}} (-1)^{\Omega(n)} = R - P = 2(R - C).$$

Theorem 6. With the cut (2.1) in force,

$$2(R - C) = \sum_{\text{endpoints of surviving cells}} (-1)^{\Omega(n)}$$

and consequently

$$R < C \iff \sum (-1)^{\Omega(n)} < 0.$$

This is the point of the paper. The inequality $R < C$ is what every criterion in Paper IV eventually reduces to. Theorem 6 says it is *identical* to the statement that a Liouville sum over the sifted set is negative. The parity problem is therefore not an external obstacle that the framework happens to run into; **it is what the framework reduces to.**

Two remarks make the shape of this clearer.

- **$R < C$ is stronger than what is needed.** Since $C - R = T - S$, the inequality asks for $T > S$: more twins than double-semiprime cells, not merely $T > 0$. The natural weaker requirement is that the number of cells with at least one composite endpoint, namely $R - S$, be less than C ; measured, this holds comfortably even at cuts where $R > C$. But using it requires a *lower* bound on S — the count of cells with both endpoints semiprime — which is itself parity-sensitive.
- **The identity is indifferent to the window.** Nothing in Theorems 1–4 refers to the length or the position of W . Lengthening the window from squares to cubes, or to $[x, 2x]$, changes C , R and S and leaves the identity untouched. That is the structural reason none of the geometric variations in Papers I and IV can help.

The same content as a covariance, which is the sharpest form we can give it. Under the cut, classify a surviving cell by the two indicators $C_L, C_R \in \{0, 1\}$ recording whether its left and right endpoints are composite, and set $u = \mathbb{E}C_L$, $v = \mathbb{E}C_R$, $\kappa = \mathbb{E}(C_L C_R) - uv$. A four-state tally of the cells — both prime, one composite either way, both composite — together with $R/C = u + v$ gives immediately

$$\frac{T}{C} = (1 - u)(1 - v) + \kappa.$$

The independent part and the correlation, and nothing else. The Frechet inequality $\mathbb{E}(C_L C_R) \geq \max(0, u + v - 1)$ then places a floor under κ , and

$$T = 0 \iff \kappa = -(1 - u)(1 - v),$$

that is, exactly at maximal anti-correlation between the two rails.

The barrier realised on this surface. At the cut $z = Q^{2/3}$ the sifting parameter is $u_B = 3$, where Buchstab's function [5] gives the primes a share $r = 1/(1 + \log 2) = 0.59062$ of the rough endpoints and the semiprimes $s = \log 2/(1 + \log 2) = 0.40938$. Independence then predicts $R/C = 2s = 0.81877$ and $T/C = r^2 = 0.34883$. Exchanging the two masses — which is precisely the ambiguity Selberg's parity example exploits — and then forcing maximal anti-correlation gives the state $(0, s, s, 1 - 2s)$, whence $R/C = 2r = 1.18123$, κ at its floor, and $T = 0$ exactly. The displacement splits as $\Delta(R/C) = 0.36246$ against $\Delta\kappa$ -carrier = 0.01364, and $-0.36246 + 0.01364 = -0.34883 = -r^2$: the small quantity is not a shortfall to be closed but the exact overlap movement that lands the configuration on zero.

Measured. Over the 793 sectors with $101 \leq P < 6300$ at this cut, the weighted mean of κ is $+0.000048$, i.e. $+0.00005$ of the Frechet floor, and by range it reads $-0.0041, +0.0021, +0.0002, -0.0001$ — collapsing by an order of magnitude per range and symmetric about zero. Individual sectors scatter between -0.32 and $+0.27$ of the floor. At $P = 6229$: $u = 0.3983, v = 0.4185, \kappa = -0.00032, (1-u)(1-v) = 0.3499$ against $T/C = 0.3496$.

So the wall has a name and a number. Twins are plentiful for one reason, that κ sits at zero; and the parity obstruction is exactly that nothing here forbids κ from reaching its floor. The missing statement is about a **correlation between the two rails**, not about a density and not about a sieve — which is why sieve tools were always going to arrive at this point and stop.

3. Why an ordinary sieve cannot cross it

3.1 The classical loss is exactly a factor of two

One might hope the sieve could deliver $R < C$ directly: bound C from below and R from above. It cannot, and the loss is a clean constant.

Set the problem in the standard normalisation. Let the sequence be sifted to depth z with level of distribution $D = x^\theta$, so the sieve variable is $s = \log D / \log z$, and let f_1, F_1 be the linear-sieve functions, the solutions of

$$(sF)' = f(s-1), \quad (sf)' = F(s-1), \quad sF(s) = 2e^\gamma \quad (1 \leq s \leq 3), \quad f(s) = \frac{2e^\gamma \log(s-1)}{s} \quad (2 \leq s \leq 4).$$

The lower bound for the survivor count carries $f_1(s)$; the Buchstab upper bound for the composite part is, after the substitution that removes both θ and the depth,

$$I(s) = \int_1^{s-1} \frac{F_1(v)}{s-v} dv.$$

Theorem 7. $I(s) = 2f_1(s)$ for $2 \leq s \leq 4$.

*(Beyond $s = 4$ the equality fails and the excess is in our favour, but we do not prove that: computed from the delay system, $I(s)/2f_1(s) = 1.011, 1.044, 1.076, 1.079$ at $s = 5, 6, 8, 12$. **Measured, not proved**, and nothing below uses it.)*

Proof for the whole range $2 \leq s \leq 4$. The linear sieve upper function satisfies $vF_1(v) = 2e^\gamma$ throughout $1 \leq v \leq 3$, and for $s \leq 4$ the integration variable runs over $v \in [1, s-1] \subseteq [1, 3]$, so $F_1(v) = 2e^\gamma/v$ on the whole range of integration. Partial fractions then give

$$I(s) = \frac{2e^\gamma}{s} \int_1^{s-1} \left(\frac{1}{v} + \frac{1}{s-v} \right) dv = \frac{4e^\gamma}{s} \log(s-1),$$

while $f_1(s) = 2e^\gamma \log(s-1)/s$ on $2 \leq s \leq 4$. Hence $I(s) = 2f_1(s)$ there. ■

(An earlier draft proved this only on $[2, 3]$ and verified $[3, 4]$ numerically. The restriction was unnecessary: the initial range of F_1 already covers $v \leq 3$, which is all $s \leq 4$ requires.)

Verification. Computed from the delay system: $I(s)/2f_1(s) = 1.00000$ at $s = 2.2, 2.5, 3.0, 3.5, 4.0$, confirming the theorem on its whole range. Checks on the functions themselves: $f_1(3) = 0.82303 = 2e^\gamma \log 2/3$, $F_1(2) = 1.781072 = e^\gamma$, $F_1(3) = 1.187382 = 2e^\gamma/3$.

Corollary 1. On $2 \leq s \leq 4$ the naive balance is $f_1(s) - I(s) = -f_1(s)$, identically. The deficit is a **factor of two, uniform on that range and independent of θ and of the window** — and $[2, 4]$ is the whole of the sieve-usable region for this problem, since a lower bound in dimension two needs $s > \beta_2 = 4.2664$ (§3.2).

That uniformity is the reason no choice of cut, level or geometry has ever improved matters in this framework: the parameters cancel out of the comparison before the comparison is made. It is Selberg's parity factor in its most explicit form, reached here from inside the construction.

3.2 The two constraints on the cut are incompatible

The heuristic version of $R < C$ is also instructive, because it fails by a small and identifiable amount.

Among the z -rough numbers below x , the proportion that are prime is $1/(u\omega(u))$ with $u = \log x / \log z$ and ω Buchstab's function; hence

$$\frac{R}{C} = 2\left(1 - \frac{1}{u\omega(u)}\right), \quad \text{so} \quad R < C \iff u\omega(u) < 2.$$

Since $\omega(u) \rightarrow e^{-\gamma}$ rapidly, the numerical threshold is $u^* \approx 3.5658$, close to the limiting proxy $2e^\gamma \approx 3.56215$; the two are near but not equal, and we use u^* where the numerics require it and $2e^\gamma$ only as the limiting value. Equivalently: **$R < C$ asks that more than half the rough numbers be prime.**

At the natural cut the threshold is closed-form, and needs no delay system. The cut (2.1) is $z^3 > U$, i.e. $u = 3$, and there $\omega(u) = (1 + \log(u - 1))/u$ still holds, so

$$u\omega(u)|_{u=3} = 1 + \log 2, \quad \frac{R}{C} \rightarrow 2\left(1 - \frac{1}{1 + \log 2}\right) = 0.818768.$$

The same number arrives without Buchstab at all: a z -rough $n \leq x$ with $z = x^{1/3}$ is prime or a product of two primes $x^\alpha, x^{1-\alpha}$ with $\alpha \in (1/3, 1/2)$, and counting the latter gives $\frac{x}{\log x} \int_{1/3}^{1/2} \frac{d\alpha}{\alpha(1-\alpha)} = \frac{x}{\log x} \log 2$, so $\Phi(x, x^{1/3}) \sim (1 + \log 2)x / \log x$ and the prime proportion tends to $1/(1 + \log 2) = 0.590616$. **So at the cut this paper actually uses, the heuristic of this subsection is a consequence of the prime number theorem and not of any sieve estimate.**

Measured, at $M = h = 1009$ with the cut at U^α ; the predicted row is $2(1 - 1/u\omega(u))$:

α	0.20	0.25	0.28	1/3	0.38	0.42
u	5.08	4.02	3.58	3.00	2.63	2.38
R/C , measured	1.2025	1.0369	0.9453	0.7559	0.5806	0.4003
R/C , predicted	1.2988	1.1139	1.0040	0.8188	0.6564	0.4872

(The prediction row is now complete and recomputed at the same u from the delay system; the $\alpha = 0.25$ entry was printed as 1.110 in an earlier version and is 1.1139.)

The prediction sits systematically above the measurement, and the gap is a finite-size effect that closes exactly. Taking the global window — every cell $6n \pm 1 \leq X$, cut at $z = X^{1/3}$, so $u = 3$ throughout — and classifying each surviving endpoint by Theorem 1, with Φ written for $\Phi(x, x^{1/3})$ in the last row:

X	10^6	10^7	10^8	10^9	4×10^9	limit
C	19,303	142,921	1,096,286	8,775,268	30,857,268	—
T	8,168	58,979	440,311	3,424,505	11,944,437	—
R/C	0.7038	0.7208	0.7355	0.7527	0.7575	0.8188
T/C	0.4231	0.4127	0.4016	0.3902	0.3871	0.3488
$\Phi \log x/x$	1.6683	1.6753	1.6783	1.6893	—	$1 + \log 2 = 1.6931$

The prime proportion among surviving endpoints is the quotient of two finite-size quantities, $[\pi(x) \log x/x]/[\Phi(x, x^{1/3}) \log x/x]$; at $X = 10^9$ that reads $1.0537/1.6893 = 0.6237$ against the directly measured 0.6237. **The whole distance from the prediction is the secondary term of the prime number theorem, and the numerator is what carries it.** The limiting value of T/C shown above is $1/(1 + \log 2)^2 = 0.3488$, i.e. the two endpoints treated as independent; the measured departure from independence, $(T \cdot S)/(\text{one-composite halves})$, reads 1.0633, 1.0694, 1.0354, 1.0237, 1.0203 across the same five heights — decreasing, consistent with independence, and **not established by it**, since that residual is the only place a twin excess could live.

On the other side, a positive lower bound for C is a sieve lower bound in **dimension two**, and therefore requires $s > \beta_2 = 4.2664$.

$$u < 2e^\gamma = 3.5621 \quad \text{and} \quad s > \beta_2 = 4.2664 \quad \text{cannot both hold.}$$

The window is empty, and it stays empty under the dimension-versus-level trade: taking $\kappa = 1$ with $\theta = 1/2$ (Chen's setting) requires $u > \beta_1/\theta = 4$, exactly what $\kappa = 2$ with $\theta = 1$ requires. **The trade between dimension and level of distribution is neutral for this problem.**

A caution about the last paragraph. This comparison uses only leading-order densities; it ignores the sieve efficiency factors f_κ , which vanish as $s \rightarrow \beta_\kappa$, and $F_\kappa > 1$ in the upper bound. Including them makes the requirement strictly harder — the tell is that the row $\kappa = 1$, $\theta = 1$ (Elliott–Halberstam) gives $u > 2$, which lies inside the window, and Elliott–Halberstam is known not to give twin primes. The comparison above is therefore a diagnostic, not a criterion, and Theorem 7 is the criterion.

3.3 Why switching cannot repair it

Chen's switching principle bridges exactly the region where the upper-bound sieve is vacuous, that is where $D/(q_1 q_2 \cdots) < z$. Its reach can be located precisely in the present normalisation.

For the target $\Omega \leq 2$ the quantity to subtract is the double Buchstab integral, and the balance $f_1(s) - I_2$ is **positive** on the sieve-usable region: computed at $\theta = 1/2$, it is +0.823, +0.789, +0.583 and +0.128 at $u = 6, 7, 8, 10$, going negative just afterwards. Chen's own choice is $z = x^{1/10}$, i.e. $u = 10$, where the balance is +0.128 — thin and positive, which is a strong check that the normalisation is right.

For the target $\Omega \leq 1$, by contrast, Corollary 1 gives $-f_1(s)$ on the sieve-usable region, **before the vacuous region is considered at all.**

The switching principle is not too weak for the twin problem; it is aimed at the wrong region. For $\Omega \leq 2$ the usable region is already positive and switching need only show the vacuous region is small — which Chen proves. For $\Omega \leq 1$ the usable region is itself short by a factor of two, and switching has nothing to repair.

3.4 Five routes through the line geometry, closed by measurement

The identity of §2.2 says what must be shown; the line geometry offers five ways of trying to show it, and all five were followed in one explicit window and closed there. They are, in order: **prime gaps** (each P_2 endpoint has a unique smallest factor, so the deficit is a statement about gaps between those factors); **line capacity** (the later lines have far more raw strikes available than a single twin needs); **the two endpoints** (whether the two members of a cell behave independently); **the two small factors** (whether the pair of factors of a P_2P_2 cell carries usable structure); and **the partner condition** (what the requirement that both members survive is worth).

The measurements are consistent with one another and with §3.2. The deviation from independence across the window is under three parts in a thousand and changes sign; the partner condition supplies a factor of about 0.17 and is the whole of the margin; and a bound that does not see both members of a cell at once fails by a factor of five before the sieve constant is reached.

The five routes, with the window they were run in and the numbers that close each, are in Appendix C.1.

3.5 The resources are ample: a control that isolates the residues

The five routes above all ask whether the lines have *enough* to close a window. This subsection answers that question separately, and the answer removes the whole family of arguments at once.

The control. Fix a window $(m^2, (m+2)^2)$ in cell coordinates and give the lines exactly the resources they actually have: one line per prime $5 \leq p \leq m$, two residue classes closed per line. Now let the classes be **chosen freely** rather than forced to $\pm 6^{-1} \pmod{p}$ — take them greedily, each line removing as many still-open cells as it can. Everything else is identical: the same lines, the same count of classes, the same window.

m	cells	lines	left open by the true classes	left open by a greedy choice
101	68	24	8	0
499	333	93	13	0
1009	673	167	27	0
2001	1335	301	50	0
4001	2668	549	70	0
10007	6672	1228	161	0
20011	13341	2261	263	0
50021	33348	5132	571	0
100003	66669	9591	995	0

With identical resources, freely allocated, the window is covered completely every time — up to 9,591 lines against 66,669 cells, where the true classes still leave 995 survivors.

What this settles. Every quantity this programme has measured on the strike side is a statement about resources: the sum $\sum 2/q$, the band decomposition of a window, the capacity of a single line, the closing budget, the layer ceilings, the six positions read as a budget. The control says that none of them can be the obstruction, because the resources are sufficient by a wide margin and the covering fails anyway. **What prevents it is only that the classes are forced to be $\pm 6^{-1} \pmod{p}$, and nothing else.** Any proof must therefore use the arithmetic of those specific residues; no argument about size, count, capacity, step or window length can reach the conclusion, because such an argument would prove the greedy case too, and the greedy case is false.

That is the same content as Theorem 6 approached from the covering side rather than the analytic one, and it is worth having in both forms. It is also the same lesson as the global maximum-gap criterion of [23]: a criterion that protects *every* translation is far stronger than one that protects the particular translations the construction actually produces, and it is the second, weaker statement that carries the problem.

3.6 The trap this section exists to avoid

Theorem 6 is worth restating as a discipline rather than only as a result. The identity

$$2(R - C) = \sum (-1)^{\Omega(n)}$$

says that the inequality one wants is *equivalent* to a statement nobody knows how to prove. It would have been easy, and would have looked like progress, to write the same content in a form that hides this.

The characteristic way of hiding it is to introduce a ratio and assume it bounded. Suppose one writes the total correlation over all shifts as a sum of the individual shift correlations,

$$\sum_{n < m \leq x} f(n)f(m) = \sum_{l \geq 1} \sum_{n \leq x} f(n)f(n+l),$$

an identity — both sides count the ordered pairs $n < m$ — and then bounds each term by a multiple of the one term of interest,

$$\sum_{n \leq x} f(n)f(n+l) \leq C \sum_{n \leq x} f(n)f(n+l_0),$$

with C declared to be a constant. Inverting gives a lower bound for the shift- l_0 correlation in terms of the total, and for $f = \vartheta$, $l_0 = 2$ the total is $\sim x^2/2$ by the prime number theorem alone, so a positive lower bound for the twin correlation appears to follow.

It does not, and the reason seems to us instructive. The quantity declared constant is

$$C = \max_{l \leq x} \frac{\sum_n f(n)f(n+l)}{\sum_n f(n)f(n+l_0)},$$

whose denominator is the very quantity being bounded. Assuming C bounded independently of x is assuming that the shift- l_0 correlation is not asymptotically smaller than every other shift correlation — which for $l_0 = 2$ and $f = \vartheta$ is the twin prime conjecture.

The step is not merely unproved; the inequality is false as a general statement about f . Take f to be the indicator of the multiples of 3 together with the two extra points 5 and 7. The shift-2 correlation is then the three pairs (3, 5), (5, 7), (7, 9) and is **equal to 3 for every x** , while the double sum divided by x grows like $x/18$:

x	10^2	10^3	10^4	10^5
shift-2 correlation	3	3	3	3
(double sum)/ x	6.0	55.9	555.9	5555.9
ratio required of C	2.0	18.6	185.3	1852.0

No fixed C exists.

We record this because the shape of the fallacy is exactly the shape of the correct statement. Every criterion in Paper IV and every measurement in §§4–7 below reduces, on inspection, to a ratio that would have to stay bounded; §§2–3 are written to name that ratio rather than to assume it. **The discipline the whole framework has tried to keep is this: when a construction reduces to an unknown quantity, the useful thing to do is to identify the quantity, not to give it a name and a positive sign.**

4. How the earlier routes reach the same wall

Sections 5 to 7 of an earlier version exhibited three routes that the framework produced before the identity of §2 was written: the **closing budget**, which misses by a conditional margin of under one per cent; a **deterministic local constraint** on how many lines must cooperate to erase a family; and the three **inheritance laws** read together. Each reaches the same place, and none crosses it.

They are worth having on record because each was pursued to its end and because the way each fails is the same way: a lower bound on the survivors is available, an upper bound on the composite part is not, and the two are needed together. That is the shape of §2.3 in three different notations.

The three routes, with their measurements and the exact point at which each stops, are in Appendix C.2.

8. Scope: two test cases, the sharpest form of the requirement, and the pattern of the whole

The examples above suggest a useful organising pattern, which we state explicitly as a summary of the tested cases rather than as a universal classification. The section then runs two test cases whose answers are known independently (§8.2–§8.4 and §8.5–§8.6), states in §8.7 the narrowest form the twin requirement takes, and closes in §8.8 with the pattern common to every route examined in this paper.

8.1 The pattern seen in the test cases

periodic statement / zero prime input	decided internally in these examples	gap 2
one prime in a progression	Dirichlet supplies the missing input here	gap 4
two simultaneous primes	not decided by the present framework	gap 6 / twins

This table summarises the examples studied in this paper; it is not asserted as a theorem classifying every prime problem. The framework describes the periodic divisibility pattern exactly, while simultaneous primality requires information not contained in those periodic counts alone.

8.2 A test case: Jacobsthal's function

Jacobsthal's function $j(n)$ [17] — the maximal gap between integers coprime to n — is such a statement: no primality enters, only the pattern modulo a primorial. Computing directly from the pattern, $j(P_k) = 2, 4, 6, 10, 14, 22, 26$ for $k = 1, \dots, 7$ (OEIS A048670 [24]) over cycles 2, 6, 30, 210, 2310, 30030, 510510.

The literature states these bounds in the **sieving bound** y rather than in $k = \pi(y)$, and the translation costs a factor of $\log$, so we give both columns. Write $P(y)$ for the product of the primes up to y , and $\log_2, \log_3$ for iterated logarithms.

	in y	in k , via $y \approx k \log k$
upper bound (Iwaniec [15], via the linear-sieve error term [14])	$j(P(y)) \ll \frac{y^2}{y^2}$	$\ll (k \log k)^2$
lower bound (Ford–Green–Konyagin–Maynard–Tao [10])	$j(P(y)) \gg y \log y \log_3 y / \log_2 y$	$\gg k \log^2 k \log_3 k / \log_2 k$
conjecture (Maier–Pomerance [21])	$j(P(y)) \ll y(\log y)^{2+o(1)}$	$\ll k(\log k)^{3+o(1)}$

Measured against the two shapes, from the tabulated values [24] for $k \leq 64$: $j/(k \log^2 k) = 1.081, 0.868, 0.969, 0.951, 0.988, 0.996$ at $k = 5, 10, 20, 30, 40, 50$ — flat at 1; while $j/(k \log k)^2 = 0.216, 0.087, 0.048, 0.032, 0.025, 0.020$ — falling, and falling *exactly* like $1/k$, which is the ratio of the two shapes.

Two readings have to be kept apart, and we separate them because it is easy not to. (i) The data sit at $k \log^2 k$, that is, at the FGKMT lower bound up to the factor $\log_3 k / \log_2 k$, and about a factor $\log k$ *below* the Maier–Pomerance conjecture. So the numbers do not confirm that conjecture; they merely fail to contradict it. (ii) The proved upper bound is away from the data by a factor of order k , not of order $\log^2 k$ — which is exactly what the clean $1/k$ decay of the second row records.

8.3 Proposition 2 (saturation of extremal runs)

Applying the framework to the *structure* rather than the *size* of an extremal configuration reveals a clean measured pattern.

Proposition 2 (measured). In the maximal closed run of length L for the first k primes, every line p covers exactly $\lceil L/p \rceil$ positions, with the single exception of the largest line, which covers exactly one.

primes	L	measured coverage	ceilings $\lceil L/p \rceil$
$\{2, 3\}$	3	2, 1	2, 1
$\{2, 3, 5\}$	5	3, 2, 1	3, 2, 1
$\{2, 3, 5, 7\}$	9	5, 3, 2, 1	5, 3, 2, 2
$\{2, 3, 5, 7, 11\}$	13	7, 5, 3, 2, 1	7, 5, 3, 2, 2
$\{2, 3, 5, 7, 11, 13\}$	21	11, 7, 5, 3, 2, 1	11, 7, 5, 3, 2, 2

In every displayed extremal run, no line wastes a strike relative to the stated ceiling pattern. The proportion of *critical* positions — those covered by exactly one line — rises steadily: $3/3$, $4/5$, $7/9$, $10/13$, $16/21$.

8.4 And why it yields no bound

The counting condition that saturation feeds is $\sum_p \lceil L/p \rceil \geq L$. Because $\sum_p 1/p$ diverges, this is satisfied for L into the tens of thousands from $k = 3$ onward:

k	3	5	10
true $j(P_k) - 1$	5	13	45
counting bound	$> 2 \times 10^4$	$> 2 \times 10^4$	$> 2 \times 10^4$

The bound exceeds the truth by factors of 400 to 4000 already at small k . **This is the same failure recorded in App. C.2.2 and C.2.1: the framework’s local statements are exact, and its summation over lines destroys them.**

8.5 A second test case: almost-primes between squares

The pattern in §8.1 suggests a second family of test problems, structurally closer than Jacobsthal: statements of the form *every interval $(n^2, (n+1)^2)$ contains an integer with $\Omega(m) \leq k$* . These are natural here because Ω is **precisely inheritance depth**: a point of multiplicative depth k is a point at which k layers of inheritance met.

The literature. Brun [4] obtained $k = 11$ for large n ; Chen [7] later obtained $k = 2$ for sufficiently large n . Explicit statements valid for *every* n are much harder: Dudek and Johnston [9] reached $k = 4$ using Kuhn’s weights [20], and Campbell [6] reached $k = 3$, combining a finite verification for $n^2 \leq 10^{31}$ with Richert’s logarithmic weights [27] and the explicit linear sieve of Bordignon, Johnston and Starichkova [1]. Between consecutive **cubes**, Johnston, Sorenson, Thomas and Webster [18] obtain $k = 2$ for all n . Legendre’s conjecture, $k = 1$, remains open even under the Riemann hypothesis [25].

The framework reproduces the exponent-level squares-versus-cubes contrast. A survivor of sieving to depth x^θ near x has $\Omega \leq \lfloor 1/\theta \rfloor$, while the window enters through the sieve variable $s = \alpha/\theta$ when its length is x^α :

window	$\theta = 1/4$	$\theta = 1/3$	$\theta = 1/2$
squares, $\alpha = 1/2$	2.000	1.500	1.000
cubes, $\alpha = 2/3$	2.667	2.000	1.333

At the depth required for $\Omega \leq 2$ the cube window gives a sieve variable larger by exactly the exponent gap $2/3 - 1/2 = 1/6$, i.e. by 33%. **This exponent gap is one structural reason the cube problem gives more sieve room than the square problem; it is not, by itself, a proof of the published distinction.** (Measured, the distinction is about provability and not truth: the minimum of Ω over the window is 1 for both squares and cubes at $x = 10^6, 10^8, 10^{10}$.)

Where the framework stops, and where it does not. For $k = 3$ the sieve variable is $s = 1.500$, and for $k = 2$ it is $s = 1.000$ — both far below the linear sieve’s lower-bound threshold $s > 2$. The published proofs bridge the gap with **weights**. Paper II shows that the framework carries those weights exactly on the cycle, including Richert’s once the generating function is refined by the size of the line [II, Cor 2]; Paper III finds that weights of that soft shape transfer very accurately on the tested windows [III, Prop 1]. **Thus the experiments do not identify the algebraic representation of the weight as the bottleneck; the required rigorous lower-bound argument remains external.**

How much is lost in that transfer, measured. The statement just made is qualitative, and it can be replaced by numbers. Compare the cycle law of [II, Thm 2] against a direct census over consecutive square windows $[M^2, (M + 2)^2]$ near X , sieving at depth $z = X^{1/u}$, and separate the comparison into three layers.

Layer 1 — the mean is exact. The mean inheritance depth is $\sum_{q \leq z} 2/q$, and each line strikes exactly $2/q$ of the cells of any interval up to a bounded edge term; measured, the window/cycle ratio is 1.0000 at $u = 2, 3, 4$. There is no transfer problem at this level.

Layer 2 — the state totals move by a known constant. At $u = 2$ the NN total in the window, divided by its cycle value, converges to the two-dimensional Buchstab factor $e^{2\gamma}/4 = 0.793055$ already met in Appendix B:

X	10^6	10^7	10^8	10^9	10^{10}
NN window / NN cycle	1.0113	0.9814	0.8653	0.8071	0.7953
distance to $e^{2\gamma}/4$	+0.218	+0.188	+0.072	+0.014	+0.002

Three decimal places at $X = 10^{10}$. So this layer is not a free parameter: it is a constant the framework already computes.

Layer 3 — the shape of the depth distribution carries a residual that does not decay. Within the state OO , the conditional mean depth in the window is below its cycle value by +0.18 per cent, −0.24 per cent, −1.51 per cent, −1.89 per cent, −1.78% at the same five heights, and the total-variation distance of the conditional law is 0.012, 0.028, 0.042, 0.046, 0.045. **Both saturate; neither is a finite-size effect.** Rescaling each state to its measured total removes almost none of it ($0.0770 \rightarrow 0.0755$ at $X = 10^8$), so the correction is not a single multiplicative constant. Across u the residual is signed: negative on $(2, 2.7)$ with a minimum near $u = 2.2$, positive

on $(2.7, 3.3)$, and indistinguishable from zero for $u \geq 3.5$ — a shape consistent with a Buchstab-type delay structure in the depth variable, though we do not identify it.

And the quantity the proofs actually use transfers far better than the law it is computed from. Taking Richert’s weight itself, with the sifting range and the weight range both scaled with X ($z = X^{1/4}$, $y = X^{1/2}$), and comparing the cycle prediction with the direct window census:

X	z	y	mean w , window / cycle	$\sum \max(0, w)$, window / cycle
10^8	100	10^4	0.99860	0.99610
10^9	177	3.2×10^4	0.99886	0.99628
10^{10}	316	10^5	0.99892	0.99638

The truncated sum — the quantity a weighted sieve argument evaluates — transfers with a **relative error of 0.36%, stable across three decades**; the linear part is off by 0.11%, and the truncation triples the discrepancy, as one expects of a nonlinear functional. This is five times smaller than the 1.8% distortion of the depth law from which the weight is computed, and it is the same phenomenon as [III, Prop 1]: soft weights transfer, sharp indicators do not, and OO is a sharp condition while w is a soft one.

What this does and does not settle. A weighted-sieve argument needs the *sign* of the truncated sum, not its value to four figures; an error of a third of a percent cannot change a sign. **So the computational side of the weight is not the obstruction, by a wide margin — and that is now a measured statement rather than an impression.** What is missing is not a better computation but a *proof*: a rigorous upper bound for $|\sum_{\text{window}} \max(0, w) - \sum_{\text{cycle}} \max(0, w)|$, valid for all X rather than observed over three decades. That is the distribution of a weighted Ω in a short interval, and it is the external ingredient of §9.1 in its most concrete form the present work can give it.

A sharper localisation of what is missing, obtained inside the framework. The bound above ($s = 1.500$ for $k = 3$) describes the sieve at the depth that makes survivors P_3 directly. One may instead sieve *less* deeply and let the almost-prime bound come from the size of the window. Write $Y = \sqrt{n+1}$ and sieve $(n^2, (n+1)^2)$ only to depth $Y/2$. Every prime factor of a survivor then exceeds $Y/2$ while the survivor itself is below Y^4 , and $(Y/2)^5 > Y^4$ as soon as $Y > 32$; hence

$$\Omega(x) \leq 4 \quad \text{for every survivor, once } Y > 32,$$

so the *only* obstruction to $\Omega \leq 3$ in the window is a survivor with exactly four prime factors. Such a survivor is severely constrained. Writing $x = pqrs$ with $p \leq q \leq r \leq s$:

- **the fourth factor is forced.** Given p, q, r , the admissible s lies in a window of width $(2n+1)/pqr < 2Y^2/(Y/2)^3 = 16/Y$, which is below 1 once $Y > 16$: there is at most one candidate;
- **the first two are confined to about four positions.** Given r, s , the product pq lies in a window of width $(2n+1)/rs < 2Y^2/(Y/2)^2 = 8$;
- **the factorisation always straddles n :** $pq \leq n < rs$ in every case examined (78 of 78 at $Y = 29, 53, 101, 199, 293, 401$).

Y	29	53	101	199	293	401
survivors	321	922	2,840	9,535	19,093	33,379
with $\Omega = 4$	1	3	6	15	25	28
fraction with $\Omega \leq 3$	0.99688	0.99675	0.99789	0.99843	0.99869	0.99916
bound $16/Y$ (widest s -window)	0.552	0.302	0.158	0.080	0.055	0.040
widest s -window measured	0.245	0.202	0.083	0.040	0.036	0.023

And this is what makes the localisation informative rather than encouraging. Since the exceptional set is so rigidly determined, it is natural to look for an injection sending each $\Omega = 4$ survivor to a nearby $\Omega \leq 3$ one; four constructions were tried (sliding the large half, reflection in the window, replacing the least factor by a neighbouring prime, and perturbing the difference-of-squares representation), and the best succeeded on 12 of 61 exceptional survivors — none is close to the required 100%. The measurement above explains why an injection was the wrong instrument: **the exceptional set does not appear to be the difficulty.** At $Y = 401$ it is 0.08% of the survivors and falling, so essentially the entire content of the statement is a positive *lower bound* on the number of survivors of sieving to depth $Y/2$ in a window of length $2n + 1$. That is a dimension-one sieve question with sieve variable

$$s = \frac{\log(2Y^2)}{\log(Y/2)} \rightarrow 2,$$

i.e. the critical point once more. **The reduction is genuine and sharp — it removes every degree of freedom but one — and the one it leaves is exactly the lower bound the framework cannot supply, at exactly the value of s at which it never can.** This is §8.1's pattern in its most compressed form: the framework can localise the difficulty to a set of density 10^{-3} and still not cross it.

8.6 The correlation of the survivor count, and the order of its sum

The variance of the survivor count, and where that question already stands. The measurements above concern the mean; one may also ask for the second moment of $B(a)$, the number of cells of $(a^2, (a+2)^2)$ surviving all lines $p \leq a/2$. Expanding the square introduces the correlation of the survivor indicator at cell-distance h , and **that correlation is already transported exactly by [II, Thm 4]**; nothing new is needed here beyond a change of coordinates, which we record because the resulting form is the one the present section uses.

Paper II indexes gap-2 pairs by x , so that consecutive indices differ by 2 in the integer, and obtains the ladder $K_q(h_{\text{II}}) = q - 2, q - 3, q - 4$ according as $h_{\text{II}} \equiv 0, \pm 1$, or otherwise modulo q . Cells are spaced 6 apart, so a cell-distance h is $h_{\text{II}} = 3h$; substituting, $h_{\text{II}} \equiv 0$ becomes $p \mid h$ for $p > 3$, and $h_{\text{II}} \equiv \pm 1$ becomes $3h \equiv \pm 1$, i.e. $p \mid 9h^2 - 1$. Writing $\nu_p(h) = p - K_p$ for the number of residues the pair of cells forbids:

$$\nu_p(h) = \begin{cases} 2, & p \mid h, \\ 3, & p \mid 9h^2 - 1, \\ 4, & \text{otherwise,} \end{cases}$$

the three cases being mutually exclusive, since $p \mid h$ forces $3h \equiv 0$. (Directly: with $c_p \equiv 6^{-1}$, the forbidden residues for m are $\pm c_p$ and $-h \pm c_p$, which coincide in pairs under exactly those two conditions. Checked against a direct count with no mismatch for all $p < 200$, $h < 300$.) **[II, §6.3] already names the resulting product as the singular series of the Hardy–Littlewood k -tuple conjecture, “suitably normalised”; what is added below is the normalising constant, the exact mean, and the order of the error.** The series is

$$S_B(h; z) = \prod_{5 \leq p \leq z} \frac{1 - \nu_p(h)/p}{(1 - 2/p)^2} = G(z) \prod_{p \mid h} \frac{1 - 2/p}{1 - 4/p} \prod_{p \nmid 9h^2 - 1} \frac{1 - 3/p}{1 - 4/p}, \quad G(z) = \prod_{5 \leq p \leq z} \frac{1 - 4/p}{(1 - 2/p)^2},$$

with G convergent (log of the factor is $\asymp -4/p^2$) and $G(\infty) = 0.39688$ to five places.

This series is self-normalising, exactly. Two derivations agree. Per prime, the three cases have densities $1/p$, $2/p$ and $1 - 3/p$, so with $q = 1/p$ the expected local factor is

$$(1 - 3q) + q \frac{1 - 2q}{1 - 4q} + 2q \frac{1 - 3q}{1 - 4q} = \frac{1 - 4q + 4q^2}{1 - 4q} = \frac{(1 - 2/p)^2}{1 - 4/p},$$

which is the reciprocal of the corresponding factor of G ; hence $G \prod_p E_p = 1$. Equivalently, expanding by Dirichlet convolution with $f(p) = 2/(p - 4)$ and $g(p) = 1/(p - 4)$, the condition $e \mid 9h^2 - 1$ contributes $2^{\omega(e)}$ residues and the main term is $GH \prod_p (1 + 4/(p(p - 4))) = GH \prod_p (p - 2)^2 / (p(p - 4)) = H$. This is the analogue, for the pattern $h(3h - 1)(3h + 1)$, of the singular-series averages of Montgomery and Soundararajan, and it identifies the object: S_B is the Hardy–Littlewood series of a prime quadruple, normalised by the square of the twin constant. Writing $\mathcal{H}_h = \{-1, 1, 6h - 1, 6h + 1\}$, the local factors at 2 and 3 are 8 and $27/16$, and $\prod_{p > 3}$ of the C_2 factor is $\frac{4}{3}C_2$, so that

$$\mathfrak{S}(\mathcal{H}_h) = \frac{27}{2} \left(\frac{4}{3}\right)^2 C_2^2 S_B(h) = 24 C_2^2 S_B(h).$$

The size of the error in that average is larger than a first look suggests, and we record the correction. An earlier draft of this section fitted $\sum_{h \leq H} (S_B(h) - 1)$ over $H \leq 4 \times 10^5$, obtained a coefficient near 0.85 that drifted with the fitting range, and reported the error as $O(\log H)$. Extending the computation to $H = Q_7 = 37,182,145$ (the seventh primorial from 5) shows that reading to be wrong: the drift was not instability in a coefficient but the curvature of a quadratic seen over too short a range. For the smoothed sum, which has far smaller fluctuation,

$$\Sigma_B(C) := \sum_{h < C} \left(1 - \frac{h}{C}\right) (S_B(h) - 1) = -0.0329(\log C)^2 - 0.244 \log C - 0.38,$$

with root-mean-square residual 0.0018 over nine points spanning $10^5 \leq C \leq 3.7 \times 10^7$; the linear model $a \log C + b$ has residual 0.105 over the same points, and $a \log C \log \log C + \dots$ has 0.0045. The unsmoothed sum fluctuates by $O(1)$ and does not separate the models on its own, but its local slope over the upper range is near 1.5, against the 0.85 that the short-range fit had suggested. **We therefore state the order as $(\log C)^2$ and withdraw the earlier $O(\log H)$.** The coefficient is measured, not derived; we have no closed form for it.

One consequence is worth stating, because it converts a measured constant into an explained one. In the independent-window model the dispersion is

$$\frac{\text{Var}(B)}{\mathbb{E}B} = 1 - \rho_1 + 2\rho_1\Sigma_B(C), \quad \rho_1(z) = \prod_{5 \leq p \leq z} \left(1 - \frac{2}{p}\right) \sim \frac{12C_2 e^{-2\gamma}}{\log^2 z} = \frac{2.497}{\log^2 z},$$

the factor 2 accounting for $\pm h$. Had Σ_B been $O(\log C)$ the product $\rho_1\Sigma_B$ would vanish and the dispersion would tend to 1; with $\Sigma_B \asymp -c(\log C)^2$ the two logarithms cancel and the leading term contributes $-2 \times 2.497 \times 0.0329 = -0.164$, a constant. **So the sub-Poisson behaviour is a genuine limit and not a slowly vanishing correction, and the quadratic order is exactly what makes it one.**

At the scales where the dispersion was measured, however, the subleading terms of Σ_B dominate that limit, and the model is therefore a prediction of drift rather than of a fixed number. Evaluating the fitted Σ_B at $z = a/2$, $C = 2a/3$:

a	10^4	10^5	10^6	10^8	$\rightarrow \infty$
ρ_1	0.0344	0.0213	0.0145	0.0079	0
model $\text{Var}/\mathbb{E}$	0.616	0.674	0.708	0.747	0.836
measured	0.685	0.702	—	—	—

A direct evaluation at $a = 100001$ from the exact truncated series ($C = 66,668$, $z = 50,000$, $\rho_1 = 2.13 \times 10^{-2}$, $\sum_{h \neq 0} = -13.95$) gives 0.681 against the measured 0.702, a 3% agreement with no free parameter. **The negative sign is the substance: the repulsion encoded in ν_p is what makes the count sub-Poisson.** The model runs a few percent low at these sizes, and predicts a slow rise of the dispersion toward 0.836; the two available measurements do rise, 0.685 at $a \sim 10^4$ to 0.702 at $a \sim 10^5$, against a predicted rise of 0.616 to 0.674. That the increments agree better than the levels is consistent with a constant offset from the Buchstab correction, which the independence model omits; we have not isolated it, and the rise itself remains a prediction rather than a confirmed law.

Why this is a description and not a route. It is natural to hope that a second-moment bound of the shape $\sum_{A < a \leq 2A} (B(a) - \mu(a))^2 \ll \sum_{A < a \leq 2A} \mu(a)$ would give $B(a) > 0$ for almost all a , which is far weaker than a large-deviation estimate. But that conclusion is already known, in a much stronger form and by different means. Bazzanella [2] proves **unconditionally** that all intervals $[n^2, (n+1)^2] \subset [1, N]$ with at most $O(N^{1/4+\varepsilon})$ exceptions contain the *expected number of primes* — not merely a survivor of the small lines — with $O(f(N)\log^2 N)$ exceptions on the Riemann hypothesis, $O(N^\varepsilon)$ under the Lindelöf or Density hypotheses [3], and none at all under a strong form of Montgomery’s conjecture. The mechanism is zero-density estimates for ζ together with a fourth-power mean of ψ in short intervals; the bridge to square windows is the observation that consecutive square windows are spaced further apart than the block of exceptions forced around any one of them, so their exceptional blocks are disjoint [3, Lem. 2.2]. **We record the local law above as a description of the dispersion, therefore, and not as the first step of a route: the almost-all statement it was reaching for is settled, more strongly, by analytic methods this framework does not contain.**

A methodological caution that cost two false readings here. The subtracted term $W(C)e_k(\{\lambda_p\})$ is of size C , so an unclosed tail in the expectation is amplified by C : two computations

differing only in the prime bound used for G and e_k disagreed by 0.1 at $C = 1.6 \times 10^6$, an amount larger than several of the effects under discussion, and the disagreement was traced exactly to $4/(P \log P)$ summed over the missing primes. Any reproduction must close that tail analytically or fix the bound; the accompanying script does the former.

8.7 The narrowest form of the requirement, and why it is still the wall

The narrowest the requirement gets, and what that shows. [IV, §3.2] reduces a twin in a sector to a single open cell avoiding six named positions. Two further facts sharpen that as far as the construction can. First, the lines 5 and 7 alone cut the six to three, and no further small line lowers it — a larger line can pick a phase striking none of the six, and the Chinese remainder theorem combines that freely with the phase of 5 and 7 that leaves three. Second, three open positions occur **only** when $(M + 2, M + 4)$ is itself a twin [IV, Cor. 2], so under the negation the count is at most two. And in five residue classes,

$$M \equiv 3, 51, 141, 153, 201 \pmod{210},$$

those two lines close all six, so there **every** open cell is a twin: the requirement becomes the bare $C_M \geq 1$.

That is the weakest *numerical* requirement the twin statement takes anywhere in this work — condition (3.2) of [IV, §3.2] is weaker still, but it asks for no count at all — and three measurements say the reduction is not paid for elsewhere. Over the 713 such sectors below $M = 30,000$ there are 444,958 open cells and **not one fails to be a twin**, against 155 non-twin survivors in the 428 ordinary sectors below $M = 3000$ alone. The survivor count itself is not depressed: the density in these five classes against all others is 0.934 up to $M = 4000$, 0.997 up to 12000 and 1.000 up to 30000 — **fixing the phases of 5 and 7 changes which cells die, not how many**.

That last sentence names a standard device rather than an observation of ours. Passing to a residue class modulo a primorial to remove the local bias at the small primes, at no asymptotic cost, is the ***W-trick***, and the modulus here is $210 = 2 \cdot 3 \cdot 5 \cdot 7$. In the analytic literature the compensating factor $\phi(W)/W$ is inserted by hand and the absence of cost is known; here it is not assumed but measured, and the three ratios above are that measurement. The device is standard — see its use in Tao and Teräväinen [30], and in the quantitative polynomial Szemerédi work of Krause, Mousavi, Tao and Teräväinen [19] — and we record the identification because the same thing has happened repeatedly here: an object built from inside the construction turns out to have a name. And C_M is nowhere near zero: its minimum over those sectors is 10 below $M = 3000$ and 520 for M near 2×10^4 , against 497 in the ordinary classes.

That calculation has been carried out, and on our reading it closes the local route. Coupling consecutive sectors and admitting the lines 11, 13, 17 bounds the exception budget over a full 210-window by 31 under the twinless hypothesis, and that ceiling cannot be lowered by any finite set of lines: of the 28 maximal configurations, 27 are killed by a fixed prime divisor once the quadratic partner conditions are included, and the survivor is admissible at every prime, so the residues can be chosen simultaneously against any finite list — the setting of Schinzel’s Hypothesis H [28]. The order in which the lines are born adds nothing either. The calculation, with its verifications, is [IV, §3.3]; here we record only what it means.

The upper-bound side of the argument is therefore finished, and the whole difficulty sits in the other half. One does not need to kill the escapee: if $\sum_i C_{M_i} \geq 32$ can be established, the thirty-second survivor has no room inside E and is a twin. Measured, that requirement is met

with enormous room — $B_k := \sum_i C_{M_i}$ is 1,337 at $k = 1$, 29,579 at $k = 100$ and 196,948 at $k = 1000$, with $B_k \log^2 M_0 / M_0$ flat at about 139 from $k = 20$ onward, and the minimum over a band of sixty consecutive k never below 1,337. **The requirement is exceeded by a factor of six thousand and rising, and remains unproved.**

And lengthening the window does not help. Sieving an interval of length H to depth z has $s = \log H / \log z$; for one sector $H \asymp 12M_0$ and for the full window $H \asymp 420M_0$, both against $z \asymp M_0$, so $s \rightarrow 1$ in each and the gain is $\Delta s = \log 35 / \log M_0$ — 0.257 at $M_0 = 10^6$, 0.129 at 10^{12} , vanishing. Against the two-dimensional sifting limit $\beta_2 = 4.2664$ this closes under seven per cent of the gap, and less as M_0 grows: **a constant multiple of the length buys nothing asymptotically, only a power of M_0 would.** The reorganisation is therefore a clarification and not a reduction — but it is a real one, because it closes one of the two branches permanently and by proof.

So the requirement is now one survivor, in one arithmetic progression, for infinitely many M — and it is still the same wall. The quantity to be bounded below is the number of cells of an interval of length $\asymp 2M$ surviving every line $p \leq M$: a two-dimensional sieve at $u = 2$, where the sifting limit is $\beta_2 = 4.2664$ and no positive lower bound is available. Lowering the demand from $C_M \geq 7$ to $C_M \geq 3$ to $C_M \geq 1$ does not approach that bound, because the sieve supplies none of the three. **The reduction is worth recording precisely because it is so sharp and still does not bind: it isolates what the framework can do — fix the arithmetic of the exceptional positions completely — from what it cannot, which is to produce a single guaranteed survivor.**

A remark from the same literature, which corroborates §8. Campbell records that reaching $\Omega \leq 2$ between squares appears to lie beyond his framework precisely because it uses only Type I information, and that this is so even for very large n [6]; Chen’s asymptotic $k = 2$ [7] comes from a different route. The deficit named *bilinearity* in §8.1 is therefore the one the specialists name as well, in the same test problem.

What does remain is the numerical optimisation itself, and the fact that at $k = 2$ even that is not enough: Campbell records that reaching $\Omega \leq 2$ appears to lie beyond an argument using only Type I information, however large n is taken. That is the fourth deficit of §8.1, met in a test problem where the first three do not bite.

8.8 The pattern of the whole work

Across the tested routes, the local laws are much sharper than the available summed bounds.

The routes examined here exhibit this repeatedly: the pigeonhole constraint (App. C.2.2), the deletion budget (App. C.2.1), the criterion $K < G$ (App. C.2.1.3), the collective cancellation [III, §4], and the covering condition for Jacobsthal (§8.4). In each case a sharp local law is followed by a summation — or, equivalently, a passage from cycle to short window — whose loss exceeds the margin available. **[II, §5.2] isolates why with unusual clarity: making the layers disjoint removes the double-counting entirely, so the summed bound stops diverging and stays below 1 for every depth — yet it then returns exactly $\prod (1 - 1/q)$.** The interference was never in the intersections; it lives in the size of each layer, which already contains the product of all the layers before it. We therefore treat this as the recurring limitation revealed by the present tests, rather than as a theorem excluding every possible summation method.

9. Conclusion

9.1 What an external ingredient would have to supply

Taking §8.5 and §8.8 together, the deficits exposed by these tests can be organised into four categories, each corresponding to a standard class of tools.

deficit	what is missing	standard remedy
Summation (§8.8)	control of the aggregate when local laws are summed over lines	large-sieve inequalities; Buchstab iteration with sign selection (Rosser–Iwaniec [16], Harman [13])
Factor size (§8.5)	a weight sensitive to <i>how large</i> the factors are, not only to how many there are	Richert’s logarithmic weights; internally, the per-bin refinement of [II, §4.3] — carried out there, and it closes
Window transfer ([III, §6])	passage from a cycle statement to a short window at $s \approx 2$	Buchstab’s function; the beta-sieve; Richert weights <i>in application</i>
Bilinearity ([III, §5.4], §8.2)	information about n as a <i>product</i> , with independent weights on the factors	Type II sums; bilinear forms; dispersion

Within this analysis the four deficits are progressively stronger. The first three concern aggregation, weighting and transfer of information already encoded by the sieve pattern. The fourth asks for bilinear information that the present framework does not generate. Papers II and III address the factor-size bookkeeping and measure the window transfer; they do not supply the missing bilinear control needed for a twin-prime conclusion.

One deficit has been narrowed rather than removed. We initially listed *continuity* — the absence of any measure sensitive to the factorisation beyond a count — as a flat impossibility. [II, Thm 2] disposes of the nonlinearity, and [II, §4.1] shows what is left is the dependence on factor size, together with a candidate repair. The lesson is worth stating, since it is the one place in this work where a stated obstruction proved to be partly an artefact of insufficient refinement.

9.2 Why the stopping point is the right one

Because it is where genuinely new information about the survivors would have to enter, rather than another rearrangement of the same sieve identities. In standard sieve terminology, the missing kind of input is **Type II / bilinear information**: control that sees an integer through interacting factors rather than only through membership in residue classes. The present cell framework supplies no such bilinear estimate for the twin-prime correlation, so the argument stops there.

What is missing is not a cleaner formulation but an arithmetic input of a different kind.

9.3 This is not a new sieve

The main sieve-level objects used across the series have classical counterparts: the dimension-2 sifting density, the sifting limit β_κ , the sieve density $V(z)$ and Buchstab-type corrections. Correspondingly,

the obstruction located here from several independent directions — residue classes, height, square position, cell states, Ω , collective cancellation, and the interaction of the two motions — is the obstruction the literature states in one sentence: **the sieve counts rough integers and does not separate a prime from a semiprime.**

What the series adds is not a new tool but a change of type. Several quantities that are usually embedded inside sieve estimates become exact cycle identities in these coordinates. Those identities help isolate where the remaining loss enters, even though the short-window conclusions still require separate estimates. Several failures come with quantitative diagnostics: the **predicted** cover margin 1.0073 under the local-constant hypothesis (App. C.2.1.2), an effective L^2 exponent about 3.16 on the tested range [III, §4], the sifting variable $s \approx 2$ against $\beta_2 = 4.2664$ (App. C.2.1.3), the Jacobsthal counting overshoot (§8.4), and the measured transfer curve from about 1.0000 to 0.80 [III, §6].

9.4 Limitations

1. **The framework is multiplicative throughout.** A line is a set of multiples, a strike is divisibility and a meeting is a common divisor. Additive prime correlations such as fixed differences are not determined by these multiplicative identities alone; an additional arithmetic input is required.
2. **Exactness on the cycle, not automatically on the window.** Paper II is exact over a cycle of length $\prod q$. Paper III shows numerically that many soft statistics transfer well, but the short-window statements are partly empirical.
3. **The closing budget is not about twins.** The graph of [IV, App. B] joins survivors at distance 6, so its conclusion concerns prime pairs $(p, p + 6)$. The caution is stated at the head of that appendix and repeated here because the phrase “gap 6” carries two meanings in this subject.
4. **No explicit constants.** The framework produces exact identities and measured ratios, not explicit numerical bounds valid from a stated point onward. The published results on almost-primes between squares are explicit-constant work, and the framework has no machinery for it.
5. **Several statements are measured, not proved,** and are labelled as such: the stability of $T/M \approx 0.80$ [III, §3.2], the saturation law of §8.3, the transfer curve [III, §6], and — in §8.6 — the quadratic order of the singular-series sum together with its coefficient 0.0329, for which we have no closed form and no proof. The exception-budget results quoted in §8.7 are of a different kind and should not be read as measurements: the ceiling of 31, the admissibility of the surviving configuration and the emptiness of the timing table are finite exhaustive computations over stated residue classes, and each is a proof. They are carried out in [IV, §3.3]. The script `verify_exception_dichotomy.py` covers the closed forms, the $\{5, 7\}$ table and [IV, Cor. 2]; the sector-coupling budget, the admissibility of the escapee and the timing table are not yet in it, and until they are they should be treated as stated but unaudited.
6. **Priority is not claimed for any result in these papers.**
7. **Negative claims deserve the same suspicion as positive ones.** We recorded one impossibility claim — that the framework could not carry analytic weights — which proved to be an artefact of insufficient refinement; Paper II removes it in two stages. Appendix B lists five further withdrawn results. **The rate at which this structure produces plausible but spurious signals is itself among the findings.**

9.5 What this framework does not presently supply

A Type II/bilinear estimate strong enough for the twin-prime correlation. The calculations in these papers do not produce such an estimate, so further rearrangement of the same periodic identities should not be mistaken for having supplied it.

References

The companion papers are cited as [0], [I], [II], [III], [IV].

1. M. Bordignon, D. R. Johnston and V. Starichkova, *An explicit version of Chen's theorem and the linear sieve*, Int. J. Number Theory **21** (2025), 2497–2572.
2. D. Bazzanella, *Primes between consecutive squares*, Arch. Math. (Basel) **75** (2000), no. 1, 29–34.
3. D. Bazzanella, *Some conditional results on primes between consecutive squares*, Funct. Approx. Comment. Math. **45** (2011), no. 2, 255–263.
4. V. Brun, *Le crible d'Ératosthène et le théorème de Goldbach*, Skrifter utgit av Videnskapsselskapet i Kristiania I, no. 3, J. Dybwad, Kristiania, 1920.
5. A. A. Buchstab, *Asymptotic estimates of a general number-theoretic function*, Mat. Sb. **44** (1937), 1239–1246.
6. P. Campbell, *On the existence of integers with at most 3 prime factors between every pair of consecutive squares*, arXiv:2603.10356 (2026).
7. J.-R. Chen, *On the distribution of almost primes in an interval*, Scientia Sinica **18** (1975), 611–627.
8. H. G. Diamond and H. Halberstam, *A higher-dimensional sieve method*, Cambridge Tracts in Mathematics **177**, Cambridge University Press, 2008.
9. A. W. Dudek and D. R. Johnston, *Almost primes between all squares*, J. Number Theory **278** (2026), 726–745.
10. K. Ford, B. Green, S. Konyagin, J. Maynard and T. Tao, *Long gaps between primes*, J. Amer. Math. Soc. **31** (2018), 65–105.
11. C. S. Franze, *Sifting limits for the $\Lambda^2\Lambda^-$ sieve*, J. Number Theory (2011); arXiv:1012.3809.
12. J. Friedlander and H. Iwaniec, *Opera de Cribro*, AMS Colloquium Publications **57**, 2010.
13. G. Harman, *Prime-Detecting Sieves*, Princeton University Press, 2007.
14. H. Iwaniec, *On the error term in the linear sieve*, Acta Arith. **19** (1971), 1–30.
15. H. Iwaniec, *On the problem of Jacobsthal*, Demonstratio Math. **11** (1978), 225–231.
16. H. Iwaniec, *Rosser's sieve*, Acta Arith. **36** (1980), 171–202.
17. E. Jacobsthal, *Über Sequenzen ganzer Zahlen von denen keine zu n teilerfremd ist*, I–III, Norske Vid. Selsk. Forh. Trondheim **33** (1960), 117–139.

18. D. R. Johnston, J. P. Sorenson, S. N. Thomas and J. E. Webster, *Primes and almost primes between cubes*, arXiv:2601.15564 (2026).
19. B. Krause, H. Mousavi, T. Tao and J. Teräväinen, *Quantitative bounds for sets lacking polynomial progressions with shifted prime difference*, arXiv:2608.19525 (2026).
20. P. Kuhn, *Neue Abschätzungen auf Grund der Viggo Brunnschen Siebmethode*, Proc. 12th Scandinavian Math. Congress (Lund, 1953), 160–168, 1954.
21. H. Maier and C. Pomerance, *Unusually large gaps between consecutive primes*, Trans. Amer. Math. Soc. **322** (1990), 201–237.
22. J. Maynard, *Small gaps between primes*, Ann. of Math. **181** (2015), 383–413.
23. T. T. K. Nguyen, *Finite-window noncovering on primorial wheels: higher-order CRT bounds and shift correlations*, Preprints.org (2026), doi:10.20944/preprints202608.1299.v1. — *A preprint, not peer reviewed; cited as contemporaneous independent work reaching the same finite-window diagnosis from the Goldbach side.*
24. OEIS Foundation, sequence A048670 (Jacobsthal’s function at the primorials).
25. J. Pintz, *Landau’s problems on primes*, J. Théor. Nombres Bordeaux **21** (2009), 357–404.
26. D. H. J. Polymath, *Variants of the Selberg sieve, and bounded intervals containing many primes*, Res. Math. Sci. **1** (2014), Art. 12.
27. H.-E. Richert, *Selberg’s sieve with weights*, Mathematika **16** (1969), 1–22.
28. A. Schinzel and W. Sierpiński, *Sur certaines hypothèses concernant les nombres premiers*, Acta Arith. **4** (1958), 185–208; erratum, ibid. **5** (1959), 259.
29. A. Selberg, *On elementary methods in prime number theory and their limitations*, Proc. 11th Scandinavian Math. Congress, Trondheim (1949), 13–22.
30. T. Tao and J. Teräväinen, *Quantitative bounds for Gowers uniformity of the Möbius and von Mangoldt functions*, J. Eur. Math. Soc. **27** (2025), 1321–1384.
31. Y. Zhang, *Bounded gaps between primes*, Ann. of Math. **179** (2014), 1121–1174.

Appendix A — Deviation tables

T/M at moving depth: see [III, §3.2].

Contribution of lines to $\sum B_r$:

U		negative	sum	positive	sum
1,019	136	−708.3	33	+93.1	
2,039	271	−2306.2	36	+162.9	

μ_r : prediction against measurement over 1,400 sectors:

r	5	7	11	17	23	101
predicted	0.200000	−0.114286	0.041558	−0.022566	−0.013013	0.001415
measured	0.200286	−0.113959	0.027162	0.010435	−0.093274	−0.037823
z	0.04	0.01	−0.44	0.49	−1.03	−0.80

Appendix B — Note on method

A single rule was followed: **every deviation was measured against an explicit baseline before being interpreted.** Several apparent results were withdrawn under it, and they are recorded here because the rate at which this structure generates spurious signals is part of the finding.

- An apparent 50% excess of gap-6 pairs in the square window proved to be a baseline error: for a pair $(n, n + 6)$ the line 3 forbids one class, not two, since $6 \equiv 0 \pmod{3}$. With the correct baseline the window is slightly *poorer* than the naive periodic average, consistent with the Buchstab-type short-window correction discussed in Paper III.
- An apparent concentration of future composites at the endpoints of gap-6 pairs proved to arise from a wrong definition of degree; with the correct definition the concentration factor is 1.007–1.015, indistinguishable from a random control sample of the same size.
- A statistically significant slope ($z = 9.10$) linking small-factor and large-factor parity proved to be an artefact of an insufficient control: a cubic polynomial failed to absorb the size dependence. With bins of width 0.02 in log and permutation within bins, the same data give $z = -1.11$.
- An apparent square-window bias of 1.5% in the e -strip proved to be a statistic artefact: averaging per-sector *ratios* rather than pooling, with only 2–4 samples per sector.
- The asymptotic constant for τ was first derived from the full-cycle density $32C_2e^{-2\gamma} = 6.6594325$, giving an apparent surplus of 13.1%. This was wrong: sieving to depth P integers of size P^2 places one exactly at $u = 2$, where the Buchstab correction is not negligible. The candidate local scale is $16C_2e^{-\gamma} = 5.9304658$, differing from the full-cycle scale by $e^\gamma/2 = \omega(2)/e^{-\gamma}$. The measurement $T(P) \log^2 P/P^2 = 5.93072$ at $P = 9973$ strongly favours this scale on the tested range, but does not establish the asymptotic constant. **Accordingly the 0.73% surplus is a conditional prediction, not a theorem.**
- A candidate zero-crossing near $2/3$ for the deviation profile is numerically close to $\sqrt{1 - e^{-\gamma}} = 0.6622239$, which the present data favour among the tested fits; and a candidate amplitude $e^{2\gamma}/4 = 0.7930547$ must **not** be treated as independent evidence, since the profile and the final deficit are linked by the exact identity of [III, identity B] and are therefore the same measurement seen twice.
- **The same constant was later reached a third time, by a third route, and it is again not independent.** Writing δ_p for the fraction of the cells still surviving that line L_p closes, the survival product $R(P) = R_0 \prod_{5 \leq p \leq P} (1 - \delta_p)$ is an exact identity, and $\prod (1 - \delta_p) / \prod (1 - 2/p)$ was measured at 0.9436, 0.9221, 0.8996, 0.8840 for $X = 10^6 \dots 10^9$, apparently approaching $e^{2\gamma}/4$. It is the same quantity as above: δ_p is computed *from the survivors*, so the product contains the twin count by construction, and substituting the Hardy–Littlewood count in integral form reproduces the whole column to four decimals (0.8842 against 0.8840 at $X = 10^9$). **The one thing in that experiment which is not a restatement is where the deviation lives:**

$\delta_p = 2/p$ to within 10^{-3} for every $p \leq X^{1/4}$ — measured ratio 1.0000 for $u \geq 5$ — with the entire departure confined to $u \in (2, 4)$, and the partial products agreeing with $(e^{\gamma\omega(u)})^2$ there. That is the fundamental lemma of the sieve stated in the vocabulary of App. C.2.1.1, and it is the reason no redistribution of effort among the lines changes anything: below $X^{1/4}$ there is nothing to redistribute.

- **A stated obstruction proved to be an artefact.** We first concluded that the framework could not carry analytic weights, since all its objects are binary or counts and a weight would require an infinite state space. Refining the four-state law by *inheritance depth* — a quantity that changes by exactly one under a strike — makes the state space finite and the law exact ([II, Thm 2]), and arbitrary weights follow. **The negative control matters here: binning by weight *value* does not close (error up to 24%), so the refinement had to be by depth and not by weight.** The general lesson is that an impossibility claim in this setting should be tested against at least one refinement before being recorded.

Any subsequent work in this structure should begin with the baseline, not with the deviation — and should treat its own negative claims with the same suspicion as its positive ones.

Appendix C — Routes that were tried and closed

The two accounts below were in the body in an earlier version. They are collected here so that the body reads as one argument — the obstruction, then why no ordinary sieve crosses it — and kept in full because a route closed by measurement seems to us worth more written down than left to be attempted again. Each has a stub in the body stating its conclusion.

C.1 Five routes through the line geometry, closed by measurement

Sections 2.1–3.3 argue that the obstruction is the distinction between $\Omega = 1$ and $\Omega = 2$ inside the sifted set. That argument is analytic. This section reports what happens when one instead asks the line geometry itself for the missing information, in a single explicit window, and follows each of the five natural routes to the point where it stops. Every number below is **measured**, not proved.

The window. We take the phase of the surviving configuration of [IV, §3.3],

$$M_0 = 448,353, \quad X = (M_0 + 210)^2 = 201,208,764,969,$$

which carries $N = 31,392,060$ cells. The cut is the largest prime with $z^3 < X$, namely $z = 5857$, since $5857^3 = 200,921,157,793$ and $5861^3 = 201,333,092,381$. After sieving to z , Theorem 1 applies and each surviving endpoint is P or P_2 . Writing the four states of a surviving cell by the status of its two endpoints:

$$C = 1,049,024, \quad R = 857,695, \quad S = 174,791, \quad T = 366,120,$$

in the letters of Theorem 4: C the surviving cells, R the P_2 endpoints among them (a cell with both endpoints composite contributing two), S the cells with both endpoints P_2 , and T the twins. The identity $T = C - R + S$ holds exactly.

Two side measurements fix the scale. The prime share among the $2R$ endpoints is $1,240,353/2,098,048 = 0.591196$ against the limit $1/(1 + \log 2) = 0.590616$ of §3.2 — agreement to four decimals, because a short window at height Y carries no secondary term in the prime

count, which is the reading §3.2 gives of its own global table. And the survivors of **all** lines up to M_0 number 366,130, that is T plus ten composite cells whose two factors both exceed M_0 ; the same measurement at $M_0 + 510,510$ and $M_0 + 1,021,020$ gives $699,747 = 699,726 + 21$ and $1,010,762 = 1,010,734 + 28$. **The survivor count of a square window is the twin count plus a two-digit remainder.**

Route 1: prime gaps. Each P_2 endpoint has a unique smallest factor $q \in (z, \sqrt{X}]$. Classify q by its gap to the next prime and compare the observed count in each class against the count predicted by the cofactor interval L/q and the prime density at X/q . Over twenty-nine gap classes from 2 to 60 the ratio is flat at 0.171 to 0.176, and the largest standardised residual anywhere is 1.59 — smaller than one expects by chance from twenty-nine classes. A control is required here and it passes: the mean gap rises from 10.16 to 12.92 across octiles of q , so gap classes are confounded with q , and the same ratio binned by q instead is flat to four decimals (1.0004, 0.9995, 1.0021, 0.9969, 0.9926, 1.0052, 1.0059, 0.9943). The gap after q is a function of the small factor alone, and the count of endpoints it owns is a prime count in the interval $(\text{lo}/q, \text{hi}/q)$; nothing links the two.

Route 2: line capacity. The 36,824 lines in $(z, \sqrt{X}]$ have 25,353,670 raw strikes available inside the window, against the ceiling $C - 1 = 1,049,023$ that a single twin needs — larger by a factor of 24.2, the usual outcome. Restricting the count to strikes that land on an endpoint of one of the R rough cells gives 857,712, against $H = 857,695$: a difference of **seventeen**, the endpoints whose two prime factors both lie below $\sqrt{X}$ and which are therefore counted twice. So the restricted capacity is not an upper bound on the output; it **equals** the output. That is Theorem 2 read as a statement about capacity — a surviving composite has exactly one responsible line, so there is no slack between what the lines can do and what they do.

Two further quantities locate why no upper bound is available here. First, the cofactor interval L/q equals the sieving depth z at $q = L/z = 32,159$; above that point the interval to be sifted is shorter than the sieve limit and no sieve estimate applies at all. Of the 36,824 lines, 34,145 lie above it, and they own 509,602 of the P_2 endpoints — 59.4 per cent of H . This is the vacuous region of §3.3, and the measurement supplies the number that section states qualitatively: for $\Omega \leq 2$ Chen’s switching principle need only show the vacuous region is small, and here it is not small, it is the majority. Second, dropping the partner condition — counting rough P_2 endpoints without requiring the other member of the cell to be rough — raises the count from 857,695 to 4,997,471, which is 4.76 times the ceiling. The partner condition supplies a factor 0.1716 and it is the whole of the margin; any bound that does not see both members of the cell at once fails by a factor of five before the sieve constant is reached. A Brun–Titchmarsh bound on the cofactor, which does not see it, gives 24,119,320.

Route 3: the two endpoints. Under independence one would predict $S = C \cdot h_L h_R$ with h_L, h_R the two P_2 shares. Measured against that prediction, over the same three windows, S gives 0.99701, 0.99847, 1.00125 and T gives 0.99857, 0.99927, 1.00060. **The deviation from independence is under three parts in a thousand and changes sign between windows.** The shares themselves are stable at 0.4086 to 0.4093, against $1 - 1/(1 + \log 2) = 0.40938$.

This also settles the status of the weaker requirement recorded in §2.3. In these letters it reads $R - S < C$; but $R - S = C - T$ identically, so the requirement is $T > 0$ — it is not a weaker route to the conclusion, it is the conclusion. What S does buy is quantitative and worth recording: discarding it via $S \geq 0$ leaves an upper bound on R a tolerance of $(C - 1)/R = 1.2231$, while keeping it raises the tolerance to 1.4269 — measured as 1.2231, 1.2230, 1.2225 and 1.4269, 1.4271, 1.4273 across the three windows. Roughly half the available margin sits in S , and a lower bound on it is a lower bound on pairs $(n, n + 2)$ with all four prime factors above $X^{1/3}$, which is the same dimension

and the same sieve variable as T .

Route 4: the two small factors. Each of the 174,791 cells of type P_2P_2 carries a factorisation of both members,

$$6j - 1 = pq, \quad 6j + 1 = rs,$$

with p, r the smaller factors. The only constraint between them is $p \neq r$, which holds in all 174,791 and follows in a line: a prime dividing $6j - 1$ leaves remainder 2 in $6j + 1$. Beyond it the pair behaves as two independent draws from the marginal law $\propto 1/(p \log(X/p))$ that the line count itself produces. The correlation of $\log p$ with $\log r$ is -0.00062 against a permutation null of mean 0.00009 and standard deviation 0.00274, that is $z = -0.26$; a chi-square on an 8×8 grid of quantiles gives 56.3 on 49 degrees of freedom, with largest standardised residual 2.47 in sixty-four cells; the share with r/p in $(0.952, 1.05)$ is 0.02264 against a shuffled 0.02188, in $(0.833, 1.2)$ it is 0.08262 against 0.08309, and in $(0.5, 2)$ it is 0.29293 against 0.29270; the median of $|r - p|$ is 73,880 against 74,042. The phase inside the sector is equally flat: correlations 0.00223 and -0.00035 with $\log p$ and $\log r$, and the ten deciles of the sector hold 17,523, 17,728, 17,247, 17,534, 17,323, 17,386, 17,552, 17,416, 17,667, 17,415 cells against a flat 17,479.

Route 5: the determinant-one relation. The same cells carry an exact Bézout relation. From $rs - pq = 2$, writing $r - p = 2a$ and $q - s = 2b$,

$$as - pb = 1, \quad \text{equivalently} \quad \left| \frac{a}{p} - \frac{b}{s} \right| = \frac{1}{ps},$$

so every P_2P_2 cell is a Farey pair, and $a \equiv s^{-1} \pmod{p}$. Verified on all 174,791. A genuine finiteness statement follows: the two congruences $6j \equiv 1 \pmod{p}$ and $6j \equiv -1 \pmod{r}$ fix j modulo $pr \geq 5861^2 = 34,351,321 > N$, so **no pair (p, r) can occur twice in the window** — and indeed the 174,791 cells carry 174,791 distinct pairs.

The relation is nevertheless local. Ordering the cells by j , the natural composition law would make consecutive cells Farey neighbours, $|a_i p_{i+1} - a_{i+1} p_i| = 1$. **That holds for none of the 174,790 consecutive pairs**; the median of that quantity is 2.3×10^9 and its minimum anywhere in the window is 4604. Nor does anything telescope: $\sum a_i/p_i = 2.6170 \times 10^5$ with total variation 5.1593×10^5 , against the 1.181 that a boundary term would leave, and the same failure by five orders of magnitude for b/s and for $1/(ps)$. Reordering the cells at random gives the same zero, so the ordering by j contributes nothing. As a bipartite graph on (p, r) the configuration is likewise generic: 16,289 four-cycles against $15,830 \pm 195$ from a degree-preserving shuffle, a ratio of 1.029.

A methodological note, because this measurement was got wrong twice before it was got right. The shuffle must be constrained to simple graphs. A closed-form configuration estimate gave 64, reading the data as a 255-fold excess; an unconstrained shuffle created 115 to 156 multi-edges and gave 20,360, reading the same data as a 20 per cent deficit. Only when the shuffle respects the constraint $pr > N$ that the geometry itself imposes — the constraint proved two paragraphs above — does the null become 15,830 and the answer 1.029. **The correct null was derivable from the object under study, and neither wrong null was distinguishable from a signal by inspection.**

What the five have in common. Each route asks the line geometry for a quantity that would separate $\Omega = 1$ from $\Omega = 2$, and each returns a quantity that is either identically the output, or independent of it to the precision available. The framework describes the one-sided structure exactly — the marginal law of the small factor, the capacity of every line, the position of every strike — and describes the two-sided structure by independence. The twin conjecture is the assertion that the two-sided structure is *not* exactly independent, by an amount far below anything measurable here.

That is §3.2 in the vocabulary of the lines rather than of the sieve, and it is the reason the account stops where it does.



C.2 The routes tried before the identity was written

This appendix reports the routes this framework produced before §2 was written; each of them stops at the identity of §2.3. They are kept, in the order they were tried, because a route closed by measurement seems to us worth more written down than left to be attempted again — with the caveat that a measurement closes a route only on the range measured, and what is offered with each is our reading of why it fails. The four items below are the short ones; the three that fail in an instructive way are taken up in full afterwards — the deletion budget in C.2.1, the pigeonhole constraint in C.2.2, and the inheritance laws in §7.

(a) Capacity. The old lines *are* able to cover a window longer than the square window: at $p = 71$ a fully closed run of 67 cells exists elsewhere, against a window of 47. **The argument “the lines are too few” is therefore dead; only the phase at p^2 can protect a twin.**

(b) Budget. The per-line cap $2W/p$ is correct, but $\sum_p 2/p$ exceeds 1 by the fourth line: for $W = 100$, $40 + 29 + 18 + 15 = 102$. (The double-counting responsible is removed by the layering of [II, §5.1], and [II, §5.2] shows that the repaired sum returns the sieve product exactly — the bound stops diverging but does not improve.) The simple summation spends the entire budget on the small lines, long before reaching the large lines for which the cap is a single position. Solving the head exactly repairs this only partially: an exhaustive scan of all 5005 offsets and all lengths for the core $\{5, 7, 11, 13\}$ gives the uniform lemma

$$\left| R_0(I) - \frac{27}{91}|I| \right| \leq \frac{98}{13},$$

with both extrema attained exactly (at $|I| = 868$, offset 2069, and at $|I| = 4137$, offset 2937) — **a constant error, independent of $|I|$** . This moves the bottleneck from the fourth line to roughly the fifteenth, but the guaranteed minimum still falls to zero: for a **fixed** window of 100 cells it reaches 0 at depth 101. With the window scaled as the sector, $W = z^2/6$, the minimum instead grows and the ratio min/mean stabilises at 0.82–0.85 across five doublings of z . **The failure is therefore of the fixed window, not of the capacity argument.**

(c) Resonance alone. No correlation was measured between boundary resonance ($r \mid q^2 - p^2$) and survival.

(d) Parity. The sieve counts *rough* integers and does not separate a prime from a semiprime. We measured this directly on the relevant quantity: after conditioning on the size of the rough part in bins of width 0.02 in log, with permutation performed within bins, the slope of $\mathbf{P}(\Omega_{>P} \text{ odd})$ against $\Omega_{\leq P}$ is -0.0013 against a null of s.d. 0.0012, i.e. $z = -1.11$. Over 1.5×10^6 samples we therefore detect **no statistically significant dependence** between the two axes under this control.



C.2.1 The closing budget, and how far it misses

C.2.1.1 Δ_q : the Buchstab factor decomposed into internal quantities Periodically a line should delete the fraction $2/q$ of the surviving pairs. Write the actual deletion as

$$E_q = \frac{2}{q}T_{q^-} + \Delta_q, \quad \text{so} \quad T_q = \left(1 - \frac{2}{q}\right)T_{q^-} - \Delta_q, \quad (5.1)$$

and iterating from a core at 13,

$$T_P = T_{13} \prod_{13 < q \leq P} \left(1 - \frac{2}{q}\right) - \sum_{13 < r \leq P} \Delta_r \prod_{r < q \leq P} \left(1 - \frac{2}{q}\right). \quad (5.2)$$

Measured at $P = 499$, ratio of actual to periodic deletion, binned by q/P :

q/P	0.0–0.1	0.1–0.2	0.2–0.3	0.3–0.4	0.4–0.5	0.5–0.6	0.6–0.7	0.7–0.8	0.8–0.9	0.9–1.0
actual/periodic	1.0012	0.9699	0.9638	1.0622	1.1578	1.2152	1.2779	1.3575	1.4074	1.4706

Individually, $q = 5, 7, 11$ give exactly 1.0000; $q = 101$ gives 0.9519; $q = 499$ gives 1.5263.

At $P = 499$, the smallest tested lines are essentially periodic while many larger lines over-delete, by as much as 47%. A structural interpretation is that a small line traverses the window many times and sees it as a full cycle, whereas a large line traverses it few times, and the position of its strikes relative to a window bounded by two squares is not random.

The measurements therefore suggest an internal way to read the correction factor

$$\frac{16C_2e^{-\gamma}}{32C_2e^{-2\gamma}} = \frac{e^{\gamma}}{2} = 0.8905362 = \frac{\omega(2)}{e^{-\gamma}} :$$

as an accumulated effect of the Δ_q , negligible for the smallest lines and positive for many lines near P . This is an interpretation of the measured profile, not a derivation of the limiting constant.

Independent numerical check. $T(P) \log^2 P / P^2$ measures 6.01865, 5.97181, **5.93072**, 5.90936 at $P = 1999, 4999, 9973, 19997$, passing near $16C_2e^{-\gamma} = 5.9304658$ around $P \approx 10^4$. These finite values favour the local 5.93 scale over the full-cycle 6.66 scale on the tested range; they do not by themselves establish a limiting constant.

Turning point. The change of behaviour occurs at $q/P \approx 0.3$, which is exactly where the cofactor at the window's lower edge, P^2/q , equals $3P$ — the window's upper root.

C.2.1.2 A conditional asymptotic comparison: the 0.73% margin The available future deletions are the surviving semiprimes

$$S_P = \sum_{P < q < 3P} [\pi(9P^2/q) - \pi(q-1)],$$

each counted once by its least factor. The prime number theorem gives the formal main term

$$S_P \sim \left(\int_1^3 \left(\frac{9}{t} - t\right) dt \right) \frac{P^2}{\log^2 P} = (9 \log 3 - 4) \frac{P^2}{\log^2 P}, \quad 9 \log 3 - 4 = 5.8875106.$$

For the cover side, the data of C.2.1.1 suggest the local scale

$$\tau \approx 5.9304658 \frac{P^2}{\log^2 P}.$$

If that measured local scale is the true asymptotic main term (and if the lower-order U, Q terms remain negligible), then the predicted ratio is

$$\frac{\tau}{S} \approx \frac{5.9304658}{5.8875106} = 1.0072960, \quad \text{a predicted surplus of 0.73 per cent.}$$

Measured at finite P , the ratio is larger and decreases in the direction of the conditional prediction:

P	τ	S	τ/S
101	2,283	1,867	1.2228
499	32,550	28,186	1.1548
997	107,439	95,613	1.1237
1,999	359,498	326,120	1.1024

Why this is not a theorem. The cover criterion itself is exact: if the available deletions are fewer than the minimum vertex cover, an edge survives. What is not proved is the required asymptotic lower bound for τ in this short window. The coefficient 5.9304658 is imported into the comparison from the measured local scale of App. C.2.1.1, precisely at the difficult $s \approx 2$ regime. Thus the 0.73% number is a conditional target margin, not a proved surplus.

As a diagnostic, this number explains the numerical difficulty. If the predicted asymptotics are correct, any proof losing appreciably more than 0.73% cannot close this budget. It also explains why the two measured sides track one another so closely.

A caution, and a correction we record. One is tempted to derive τ 's constant from the full-cycle density $32C_2e^{-2\gamma} = 6.6594325$, which would give a surplus of 13.1% instead. That is wrong: sieving to depth P integers of size P^2 places one exactly at $u = 2$, where the Buchstab correction is not negligible, and the measurements of App. C.2.1.1 favour the 5.93 local scale on the tested range. This error was made and corrected in the course of the work; it is recorded in Appendix B.

The measured inheritance profile is not the dominant distortion at $P = 1999$, on the range tested. The ratios of window to global values for V/W , T/V , U/T and Q/U are 0.9534, 0.9531, 0.9532, 0.9533. Their agreement to four decimals suggests a nearly uniform scale factor across these observables. This supports, but does not prove asymptotically, the use of a common local correction in the budget comparison.

C.2.1.3 Why $K < G$ does not close For each survivor v let $\deg_6(v) \in \{0, 1, 2\}$, let $\mathcal{C}$ be the set of survivors the future lines will remove, and put $K = \sum_{v \in \mathcal{C}} \deg_6(v)$. If $K < G$ then some gap-6 pair has neither endpoint touched — i.e. a prime pair $(p, p+6)$, by the caution at the head of [IV, App. B].

Defining the concentration factor $c_p = (K/S) \div (2G/V)$:

p	101	199	499	997	1,999	4,999
c_p	1.0074	1.0143	1.0035	1.0094	1.0105	1.0057
required bound	2.54	2.89	3.05	3.23	3.44	3.74

c_p is flat at 1.01 across five doublings while the required bound widens; three independent checks (the identity $\sum \deg_6 = 2G$; a random control sample of equal size giving $c = 1.0195, 1.0043, 1.0071, 0.99966$; and a breakdown by killing line) show that the observed 1% effect is comparable to the random controls and is not resolved as a structural signal by these tests.

The criterion nonetheless fails, and for a reason that recurs. It needs an **upper** bound on K — available, since Selberg’s upper-bound sieve has no parity obstruction — and a **lower** bound on G , which is not. Inside $(P^2, 9P^2)$ sieved to depth P , the sieve variable is

$$s = \frac{\log(8P^2)}{\log P} = 2.451, 2.301, 2.226, 2.151 \quad \text{at } P = 101, 997, 9973, 10^6,$$

i.e. always ≈ 2 and falling toward it, whereas a lower bound in dimension 2 requires $s > \beta_2 = 4.2664$, the sifting limit of the Diamond–Halberstam–Richert sieve [8] (the $\Lambda^2\Lambda^-$ sieve gives 4.516 and the Rosser–Iwaniec β -sieve 4.85; see Franze [11] for the table and [12, §11.19]). Sieve methods have of course gone a long way at bounded gaps — Zhang [31], Maynard [22] and Polymath [26] — but Polymath’s §9 shows that 6 is the floor obtainable from purely sieve-theoretic considerations, by adapting Selberg’s parity argument [29]; the present framework sits inside that same family. **The route inherits the barrier rather than avoiding it — and the irony is that the crude part (K) is available while the part one expects to know (G) is not.**

C.2.2 A deterministic local constraint, and its weakness

Proposition 1. In any contiguous block of length H , each residue of a period- q ruler occurs $\lfloor H/q \rfloor$ or $\lceil H/q \rceil$ times. Hence one old line’s deviation from its cyclic mean is < 1 for a single mark and $< r$ for a run of r marks.

This is a pigeonhole statement, not a statistical one, and it is the only non-probabilistic constraint the framework produces. *Zero violations over $q \in \{5, 7, 11, 13\}$, $H \leq 40$, every starting position.*

It is nevertheless useless at the scale required. Summing over lines gives $2\pi(p)$: that is 46 at $p = 101$, 330 at $p = 997$, and 2,454 at $p = 10,007$, against a block of length ≈ 4 . Meanwhile the actual deviation is at or below noise level:

block H	measured s.d.	pure noise	ratio
4	0.169574	0.159424	1.064
64	0.033494	0.039856	0.840
1,024	0.005987	0.009964	0.601

The provable bound and the truth differ by two to three orders of magnitude. **The failure occurs at the summation over lines**, which is precisely where the union bound and the budget argument also fail.

C.2.3 Why the three laws do not close the problem They give a supply-and-demand comparison, and the comparison goes the wrong way. Take the fixed set $\{5, 7, 11\}$, so $Q = 385$, $S = 135$, and consider the sector at $M = 2319$: 4,644 cells, of which the tail of 4,620 carries exactly 1,620 cells open after 5, 7, 11 — [IV, Thm 11] confirmed.

[IV, Thm 12] says no single later line can close more than two of any family’s twelve copies, so **at least six distinct lines** must cooperate to erase one family. But the lines are not scarce:

open cells to be closed	1,620
later lines able to reach the sector	339
total available strikes, $\sum_q 2 \cdot (12Q)/q$	9,650
oversupply	6.0×

Running the lines in order leaves 157 of the 1,620 open, and those 157 are exactly the twin pairs of the sector. So the local constraint is real and is simply absorbed by the number of lines.

And no family is protected. Distributing the 157 survivors over the 135 families gives 38 families wiped out entirely, 51 with one survivor, 37 with two, 5 with three, 3 with four and 1 with five — never more than five of twelve. Against a null in which the twelve copies of each family survive independently with probability $p = 157/1620$, the binomial prediction is 39.7, 51.2, 30.2, 10.8, 2.6, 0.4. A second sector ($M = 4629$, 125 survivors) gives observed 54, 47, 25, 8, 1 against predicted 51.5, 51.7, 23.8, 6.6, 1.2.

The constraint binds one line at a time and dissolves in aggregate. Each line is restricted to one copy per family beyond the finite exceptional set, yet three hundred such lines erase the families at a rate indistinguishable from independent chance. This is the pattern of §8.8 again, met in the sharpest local form the framework has produced: an exact per-gap restriction, and no aggregate consequence.